

# 00-pytorch-fundamentals-video

October 26, 2023

## 0.1 00. PyTorch Fundamentals

Resource notebook: [https://www.learnpytorch.io/00\\_pytorch\\_fundamentals/](https://www.learnpytorch.io/00_pytorch_fundamentals/)

If you have a question: <https://github.com/mrdbourke/pytorch-deep-learning/discussions>

```
[ ]: import torch
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
print(torch.__version__)
```

1.10.0+cu111

## 0.2 Introduction to Tensors

### 0.2.1 Creating tensors

PyTorch tensors are created using `torch.Tensor()` = <https://pytorch.org/docs/stable/tensors.html>

```
[ ]: # scalar
scalar = torch.tensor(7)
scalar
```

```
[ ]: tensor(7)
```

```
[ ]: scalar.ndim
```

```
[ ]: 0
```

```
[ ]: # Get tensor back as Python int
scalar.item()
```

```
[ ]: 7
```

```
[ ]: # Vector
vector = torch.tensor([7, 7])
vector
```

```
[ ]: tensor([7, 7])
```

```
[ ]: vector.ndim
```

```
[ ]: 1
```

```
[ ]: vector.shape
```

```
[ ]: torch.Size([2])
```

```
[ ]: # MATRIX  
MATRIX = torch.tensor([[7, 8],  
                        [9, 10]])  
  
MATRIX
```

```
[ ]: tensor([[ 7,  8],  
          [ 9, 10]])
```

```
[ ]: MATRIX.ndim
```

```
[ ]: 2
```

```
[ ]: MATRIX[1]
```

```
[ ]: tensor([ 9, 10])
```

```
[ ]: MATRIX.shape
```

```
[ ]: torch.Size([2, 2])
```

```
[ ]: # TENSOR  
TENSOR = torch.tensor([[[1, 2, 3],  
                       [3, 6, 9],  
                       [2, 4, 5]]])  
  
TENSOR
```

```
[ ]: tensor([[[1, 2, 3],  
          [3, 6, 9],  
          [2, 4, 5]])])
```

```
[ ]: TENSOR.ndim
```

```
[ ]: 3
```

```
[ ]: TENSOR.shape
```

```
[ ]: torch.Size([1, 3, 3])
```

```
[ ]: TENSOR[0]
```

```
[ ]: tensor([[1, 2, 3],  
            [3, 6, 9],  
            [2, 4, 5]])
```

## 0.2.2 Random tensors

Why random tensors?

Random tensors are important because the way many neural networks learn is that they start with tensors full of random numbers and then adjust those random numbers to better represent the data.

Start with random numbers -> look at data -> update random numbers -> look at data -> update random numbers

Torch random tensors - <https://pytorch.org/docs/stable/generated/torch.rand.html>

```
[ ]: # Create a random tensor of size (3, 4)  
random_tensor = torch.rand(3, 4)  
random_tensor
```

```
[ ]: tensor([[0.4433, 0.7119, 0.4170, 0.4409],  
            [0.8014, 0.2050, 0.3547, 0.6358],  
            [0.3007, 0.1659, 0.3462, 0.7317]])
```

```
[ ]: # Create a random tensor with similar shape to an image tensor  
random_image_size_tensor = torch.rand(size=(3, 224, 224)) # height, width, ↵  
    ↪ colour channels (R, G, B)  
random_image_size_tensor.shape, random_image_size_tensor.ndim
```

```
[ ]: (torch.Size([3, 224, 224]), 3)
```

## 0.2.3 Zeros and ones

```
[ ]: # Create a tensor of all zeros  
zeros = torch.zeros(size=(3, 4))  
zeros
```

```
[ ]: tensor([[0., 0., 0., 0.],  
            [0., 0., 0., 0.],  
            [0., 0., 0., 0.]])
```

```
[ ]: # Create a tensor of all ones  
ones = torch.ones(size=(3, 4))  
ones
```

```
[ ]: tensor([[1., 1., 1., 1.],
            [1., 1., 1., 1.],
            [1., 1., 1., 1.]])
```

```
[ ]: ones.dtype
```

```
[ ]: torch.float32
```

```
[ ]: random_tensor.dtype
```

```
[ ]: torch.float32
```

## 0.2.4 Creating a range of tensors and tensors-like

```
[ ]: # Use torch.range() and get deprecated message, use torch.arange()
one_to_ten = torch.arange(start=1, end=11, step=1)
one_to_ten
```

```
[ ]: tensor([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```

```
[ ]: # Creating tensors like
ten_zeros = torch.zeros_like(input=one_to_ten)
ten_zeros
```

```
[ ]: tensor([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

## 0.2.5 Tensor datatypes

**Note:** Tensor datatypes is one of the 3 big errors you'll run into with PyTorch & deep learning:

1. Tensors not right datatype 2. Tensors not right shape 3. Tensors not on the right device

Precision in computing - [https://en.wikipedia.org/wiki/Precision\\_\(computer\\_science\)#:~:text=In%20computer%20science,precision%20is%20the%20number%20of%20bits%20used%20to%20represent%20a%20number](https://en.wikipedia.org/wiki/Precision_(computer_science)#:~:text=In%20computer%20science,precision%20is%20the%20number%20of%20bits%20used%20to%20represent%20a%20number)

```
[ ]: # Float 32 tensor
float_32_tensor = torch.tensor([3.0, 6.0, 9.0],
                                dtype=None, # what datatype is the tensor (e.g. ↵
                                ↵float32 or float16)
                                device=None, # What device is your tensor on
                                requires_grad=False) # whether or not to track ↵
                                ↵gradients with this tensors operations
float_32_tensor
```

```
[ ]: tensor([3., 6., 9.])
```

```
[ ]: float_32_tensor.dtype
```

```
[ ]: torch.float32
```

```
[ ]: float_16_tensor = float_32_tensor.type(torch.float16)
float_16_tensor
```

```
[ ]: tensor([3., 6., 9.], dtype=torch.float16)
```

```
[ ]: float_16_tensor * float_32_tensor
```

```
[ ]: tensor([ 9., 36., 81.])
```

```
[ ]: int_32_tensor = torch.tensor([3, 6, 9], dtype=torch.long)
int_32_tensor
```

```
[ ]: tensor([3, 6, 9])
```

```
[ ]: float_32_tensor * int_32_tensor
```

```
[ ]: tensor([ 9., 36., 81.])
```

### 0.2.6 Getting information from tensors (tensor attributes)

1. Tensors not right datatype - to do get datatype from a tensor, can use `tensor.dtype`
2. Tensors not right shape - to get shape from a tensor, can use `tensor.shape`
3. Tensors not on the right device - to get device from a tensor, can use `tensor.device`

```
[ ]: # Create a tensor
some_tensor = torch.rand(3, 4)
some_tensor
```

```
[ ]: tensor([[0.7151, 0.9288, 0.0464, 0.2910],
           [0.7281, 0.5272, 0.9098, 0.3145],
           [0.9641, 0.4652, 0.8553, 0.0232]])
```

```
[ ]: # Find out details about some tensor
print(some_tensor)
print(f"Datatype of tensor: {some_tensor.dtype}")
print(f"Shape of tensor: {some_tensor.shape}")
print(f"Device tensor is on: {some_tensor.device}")
```

```
tensor([[0.7151, 0.9288, 0.0464, 0.2910],
        [0.7281, 0.5272, 0.9098, 0.3145],
        [0.9641, 0.4652, 0.8553, 0.0232]])
Datatype of tensor: torch.float32
Shape of tensor: torch.Size([3, 4])
Device tensor is on: cpu
```

### 0.2.7 Manipulating Tensors (tensor operations)

Tensor operations include: \* Addition \* Subtraction \* Multiplication (element-wise) \* Division \* Matrix multiplication

```
[ ]: # Create a tensor and add 10 to it  
tensor = torch.tensor([1, 2, 3])  
tensor + 10
```

```
[ ]: tensor([11, 12, 13])
```

```
[ ]: # Multiply tensor by 10  
tensor * 10
```

```
[ ]: tensor([10, 20, 30])
```

```
[ ]: tensor
```

```
[ ]: tensor([1, 2, 3])
```

```
[ ]: # Subtract 10  
tensor - 10
```

```
[ ]: tensor([-9, -8, -7])
```

```
[ ]: # Try out PyTorch in-built functions  
torch.mul(tensor, 10)
```

```
[ ]: tensor([10, 20, 30])
```

```
[ ]: torch.add(tensor, 10)
```

```
[ ]: tensor([11, 12, 13])
```

### 0.2.8 Matrix multiplication

Two main ways of performing multiplication in neural networks and deep learning:

1. Element-wise multiplication
2. Matrix multiplication (dot product)

More information on multiplying matrices - <https://www.mathsisfun.com/algebra/matrix-multiplying.html>

There are two main rules that performing matrix multiplication needs to satisfy: 1. The **inner dimensions** must match: \* (3, 2) @ (3, 2) won't work \* (2, 3) @ (3, 2) will work \* (3, 2) @ (2, 3) will work 2. The resulting matrix has the shape of the **outer dimensions**: \* (2, 3) @ (3, 2) -> (2, 2) \* (3, 2) @ (2, 3) -> (3, 3)

```
[ ]: # Element wise multiplication
print(tensor, "*", tensor)
print(f"Equals: {tensor * tensor}")
```

```
tensor([1, 2, 3]) * tensor([1, 2, 3])
Equals: tensor([1, 4, 9])
```

```
[ ]: # Matrix multiplication
torch.matmul(tensor, tensor)
```

```
[ ]: tensor(14)
```

```
[ ]: tensor
```

```
[ ]: tensor([1, 2, 3])
```

```
[ ]: # Matrix multiplication by hand
1*1 + 2*2 + 3*3
```

```
[ ]: 14
```

```
[ ]: %%time
value = 0
for i in range(len(tensor)):
    value += tensor[i] * tensor[i]
print(value)
```

```
tensor(14)
CPU times: user 661 µs, sys: 876 µs, total: 1.54 ms
Wall time: 1.55 ms
```

```
[ ]: %%time
torch.matmul(tensor, tensor)
```

```
CPU times: user 67 µs, sys: 29 µs, total: 96 µs
Wall time: 101 µs
```

```
[ ]: tensor(14)
```

## 0.2.9 One of the most common errors in deep learning: shape errors

```
[ ]: # Shapes for matrix multiplication
tensor_A = torch.tensor([[1, 2],
                          [3, 4],
                          [5, 6]])

tensor_B = torch.tensor([[7, 10],
```

```

            [8, 11],
            [9, 12]])

# torch.mm(tensor_A, tensor_B) # torch.mm is the same as torch.matmul (it's an
↪ alias for writing less code)
torch.matmul(tensor_A, tensor_B)

```

```

-----
RuntimeError                                Traceback (most recent call last)
<ipython-input-46-281a2d72c2ec> in <module>()
      9
     10 # torch.mm(tensor_A, tensor_B) # torch.mm is the same as torch.matmul
↪ (it's an alias for writing less code)
----> 11 torch.matmul(tensor_A, tensor_B)

RuntimeError: mat1 and mat2 shapes cannot be multiplied (3x2 and 3x2)

```

```
[ ]: tensor_B.T
```

```
[ ]: tensor([[ 7,  8,  9],
            [10, 11, 12]])
```

```
[ ]: tensor_A.shape, tensor_B.shape
```

```
[ ]: (torch.Size([3, 2]), torch.Size([3, 2]))
```

To fix our tensor shape issues, we can manipulate the shape of one of our tensors using a **transpose**. A **transpose** switches the axes or dimensions of a given tensor.

```
[ ]: tensor_B, tensor_B.shape
```

```
[ ]: (tensor([[ 7, 10],
            [ 8, 11],
            [ 9, 12]]), torch.Size([3, 2]))
```

```
[ ]: tensor_B.T, tensor_B.T.shape
```

```
[ ]: (tensor([[ 7,  8,  9],
            [10, 11, 12]]), torch.Size([2, 3]))
```

```

[ ]: # The matrix multiplication operation works when tensor_B is transposed
print(f"Original shapes: tensor_A = {tensor_A.shape}, tensor_B = {tensor_B.
↪ shape}")
print(f"New shapes: tensor_A = {tensor_A.shape} (same shape as above), tensor_B.
↪ T = {tensor_B.T.shape}")

```



```
print(f"Multiplying: {tensor_A.shape} @ {tensor_B.T.shape} <- inner dimensions_
↳must match")
print("Output:\n")
output = torch.matmul(tensor_A, tensor_B.T)
print(output)
print(f"\nOutput shape: {output.shape}")
```

Original shapes: tensor\_A = torch.Size([3, 2]), tensor\_B = torch.Size([3, 2])  
 New shapes: tensor\_A = torch.Size([3, 2]) (same shape as above), tensor\_B.T = torch.Size([2, 3])

Multiplying: torch.Size([3, 2]) @ torch.Size([2, 3]) <- inner dimensions must match

Output:

```
tensor([[ 27,  30,  33],
        [ 61,  68,  75],
        [ 95, 106, 117]])
```

Output shape: torch.Size([3, 3])

### 0.3 Finding the min, max, mean, sum, etc (tensor aggregation)

```
[ ]: # Create a tensor
x = torch.arange(1, 100, 10)
x, x.dtype
```

```
[ ]: (tensor([ 1, 11, 21, 31, 41, 51, 61, 71, 81, 91]), torch.int64)
```

```
[ ]: # Find the min
torch.min(x), x.min()
```

```
[ ]: (tensor(1), tensor(1))
```

```
[ ]: # Find the max
torch.max(x), x.max()
```

```
[ ]: (tensor(91), tensor(91))
```

```
[ ]: # Find the mean - note: the torch.mean() function requires a tensor of float32_
↳datatype to work
torch.mean(x.type(torch.float32)), x.type(torch.float32).mean()
```

```
[ ]: (tensor(46.), tensor(46.))
```

```
[ ]: # Find the sum
torch.sum(x), x.sum()
```

```
[ ]: (tensor(460), tensor(460))
```

## 0.4 Finding the positional min and max

```
[ ]: x
```

```
[ ]: tensor([ 1, 11, 21, 31, 41, 51, 61, 71, 81, 91])
```

```
[ ]: # Find the position in tensor that has the minimum value with argmin() ->  
↪ returns index position of target tensor where the minimum value occurs  
x.argmin()
```

```
[ ]: tensor(0)
```

```
[ ]: x[0]
```

```
[ ]: tensor(1)
```

```
[ ]: # Find the position in tensor that has the maximum value with argmax()  
x.argmax()
```

```
[ ]: tensor(9)
```

```
[ ]: x[9]
```

```
[ ]: tensor(91)
```

## 0.5 Reshaping, stacking, squeezing and unsqueezing tensors

- Reshaping - reshapes an input tensor to a defined shape
- View - Return a view of an input tensor of certain shape but keep the same memory as the original tensor
- Stacking - combine multiple tensors on top of each other (vstack) or side by side (hstack)
- Squeeze - removes all 1 dimensions from a tensor
- Unsqueeze - add a 1 dimension to a target tensor
- Permute - Return a view of the input with dimensions permuted (swapped) in a certain way

```
[ ]: # Let's create a tensor  
import torch  
x = torch.arange(1., 10.)  
x, x.shape
```

```
[ ]: (tensor([1., 2., 3., 4., 5., 6., 7., 8., 9.]), torch.Size([9]))
```

```
[ ]: # Add an extra dimension  
x_reshaped = x.reshape(1, 9)  
x_reshaped, x_reshaped.shape
```

```
[ ]: (tensor([[1., 2., 3., 4., 5., 6., 7., 8., 9.]]), torch.Size([1, 9]))
```

```
[ ]: # Change the view
z = x.view(1, 9)
z, z.shape
```

```
[ ]: (tensor([[1., 2., 3., 4., 5., 6., 7., 8., 9.]]), torch.Size([1, 9]))
```

```
[ ]: # Changing z changes x (because a view of a tensor shares the same memory as
    ↪ the original input)
z[:, 0] = 5
z, x
```

```
[ ]: (tensor([[5., 2., 3., 4., 5., 6., 7., 8., 9.]]),
      tensor([5., 2., 3., 4., 5., 6., 7., 8., 9.]))
```

```
[ ]: # Stack tensors on top of each other
x_stacked = torch.stack([x, x, x, x], dim=0)
x_stacked
```

```
[ ]: tensor([[5., 2., 3., 4., 5., 6., 7., 8., 9.],
           [5., 2., 3., 4., 5., 6., 7., 8., 9.],
           [5., 2., 3., 4., 5., 6., 7., 8., 9.],
           [5., 2., 3., 4., 5., 6., 7., 8., 9.]])
```

```
[ ]: # torch.squeeze() - removes all single dimensions from a target tensor
print(f"Previous tensor: {x_resaped}")
print(f"Previous shape: {x_resaped.shape}")

# Remove extra dimensions from x_resaped
x_squeezed = x_resaped.squeeze()
print(f"\nNew tensor: {x_squeezed}")
print(f"New shape: {x_squeezed.shape}")
```

Previous tensor: tensor([[5., 2., 3., 4., 5., 6., 7., 8., 9.]])

Previous shape: torch.Size([1, 9])

New tensor: tensor([5., 2., 3., 4., 5., 6., 7., 8., 9.])

New shape: torch.Size([9])

```
[ ]: # torch.unsqueeze() - adds a single dimension to a target tensor at a specific
    ↪ dim (dimension)
print(f"Previous target: {x_squeezed}")
print(f"Previous shape: {x_squeezed.shape}")

# Add an extra dimension with unsqueeze
x_unsqueezed = x_squeezed.unsqueeze(dim=0)
```

```
print(f"\nNew tensor: {x_unsqueezed}")
print(f"New shape: {x_unsqueezed.shape}")
```

Previous target: tensor([5., 2., 3., 4., 5., 6., 7., 8., 9.])

Previous shape: torch.Size([9])

New tensor: tensor([[5., 2., 3., 4., 5., 6., 7., 8., 9.]])

New shape: torch.Size([1, 9])

```
[ ]: # torch.permute - rearranges the dimensions of a target tensor in a specified
      ↪order
x_original = torch.rand(size=(224, 224, 3)) # [height, width, colour_channels]

# Permute the original tensor to rearrange the axis (or dim) order
x_permuted = x_original.permute(2, 0, 1) # shifts axis 0->1, 1->2, 2->0

print(f"Previous shape: {x_original.shape}")
print(f"New shape: {x_permuted.shape}") # [colour_channels, height, width]
```

Previous shape: torch.Size([224, 224, 3])

New shape: torch.Size([3, 224, 224])

```
[ ]: x_original[0, 0, 0] = 728218
      x_original[0, 0, 0], x_permuted[0, 0, 0]
```

```
[ ]: (tensor(728218.), tensor(728218.))
```

## 0.6 Indexing (selecting data from tensors)

Indexing with PyTorch is similar to indexing with NumPy.

```
[ ]: # Create a tensor
import torch
x = torch.arange(1, 10).reshape(1, 3, 3)
x, x.shape
```

```
[ ]: (tensor([[[1, 2, 3],
               [4, 5, 6],
               [7, 8, 9]]]), torch.Size([1, 3, 3]))
```

```
[ ]: # Let's index on our new tensor
x[0]
```

```
[ ]: tensor([[1, 2, 3],
            [4, 5, 6],
            [7, 8, 9]])
```

```
[ ]: # Let's index on the middle bracket (dim=1)  
x[0][0]
```

```
[ ]: tensor([1, 2, 3])
```

```
[ ]: # Let's index on the most inner bracket (last dimension)  
x[0][1][1]
```

```
[ ]: tensor(5)
```

```
[ ]: # You can also use ":" to select "all" of a target dimension  
x[:, 0]
```

```
[ ]: tensor([[1, 2, 3]])
```

```
[ ]: # Get all values of 0th and 1st dimensions but only index 1 of 2nd dimension  
x[:, :, 1]
```

```
[ ]: tensor([[2, 5, 8]])
```

```
[ ]: # Get all values of the 0 dimension but only the 1 index value of 1st and 2nd  
      ↪ dimension  
x[:, 1, 1]
```

```
[ ]: tensor([5])
```

```
[ ]: # Get index 0 of 0th and 1st dimension and all values of 2nd dimension  
x[0, 0, :]
```

```
[ ]: tensor([1, 2, 3])
```

```
[ ]: # Index on x to return 9  
print(x[0][2][2])  
  
# Index on x to return 3, 6, 9  
print(x[:, :, 2])
```

```
tensor(9)  
tensor([[3, 6, 9]])
```

## 0.7 PyTorch tensors & NumPy

NumPy is a popular scientific Python numerical computing library.

And because of this, PyTorch has functionality to interact with it.

- Data in NumPy, want in PyTorch tensor -> `torch.from_numpy(ndarray)`
- PyTorch tensor -> NumPy -> `torch.Tensor.numpy()`

```
[ ]: # NumPy array to tensor
import torch
import numpy as np

array = np.arange(1.0, 8.0)
tensor = torch.from_numpy(array) # warning: when converting from numpy ->
    ↳pytorch, pytorch reflects numpy's default datatype of float64 unless
    ↳specified otherwise
array, tensor
```

```
[ ]: (array([1., 2., 3., 4., 5., 6., 7.]),
      tensor([1., 2., 3., 4., 5., 6., 7.], dtype=torch.float64))
```

```
[ ]: # Change the value of array, what will this do to `tensor`?
array = array + 1
array, tensor
```

```
[ ]: (array([2., 3., 4., 5., 6., 7., 8.]),
      tensor([1., 2., 3., 4., 5., 6., 7.], dtype=torch.float64))
```

```
[ ]: # Tensor to NumPy array
tensor = torch.ones(7)
numpy_tensor = tensor.numpy()
tensor, numpy_tensor
```

```
[ ]: (tensor([1., 1., 1., 1., 1., 1., 1.]),
      array([1., 1., 1., 1., 1., 1., 1.], dtype=float32))
```

```
[ ]: # Change the tesnor, what happens to `numpy_tensor`?
tensor = tensor + 1
tensor, numpy_tensor
```

```
[ ]: (tensor([2., 2., 2., 2., 2., 2., 2.]),
      array([1., 1., 1., 1., 1., 1., 1.], dtype=float32))
```

## 0.8 Reproducibility (trying to take random out of random)

In short how a neural network learns:

start with random numbers -> tensor operations -> update random numbers to try and make them better representations of the data -> again -> again -> again...

To reduce the randomness in neural networks and PyTorch comes the concept of a **random seed**.

Essentially what the random seed does is “flavour” the randomness.

```
[ ]: import torch

# Create two random tensors
```

```

random_tensor_A = torch.rand(3, 4)
random_tensor_B = torch.rand(3, 4)

print(random_tensor_A)
print(random_tensor_B)
print(random_tensor_A == random_tensor_B)

```

```

tensor([[0.3675, 0.8410, 0.0507, 0.3165],
        [0.7275, 0.9676, 0.3901, 0.8840],
        [0.5177, 0.2239, 0.4362, 0.3602]])
tensor([[0.5229, 0.6719, 0.2790, 0.8198],
        [0.6689, 0.8659, 0.7849, 0.4268],
        [0.2076, 0.8076, 0.4377, 0.2555]])
tensor([[False, False, False, False],
        [False, False, False, False],
        [False, False, False, False]])

```

[ ]: *# Let's make some random but reproducible tensors*

```

import torch

# Set the random seed
RANDOM_SEED = 42
torch.manual_seed(RANDOM_SEED)
random_tensor_C = torch.rand(3, 4)

torch.manual_seed(RANDOM_SEED)
random_tensor_D = torch.rand(3, 4)

print(random_tensor_C)
print(random_tensor_D)
print(random_tensor_C == random_tensor_D)

```

```

tensor([[0.8823, 0.9150, 0.3829, 0.9593],
        [0.3904, 0.6009, 0.2566, 0.7936],
        [0.9408, 0.1332, 0.9346, 0.5936]])
tensor([[0.8823, 0.9150, 0.3829, 0.9593],
        [0.3904, 0.6009, 0.2566, 0.7936],
        [0.9408, 0.1332, 0.9346, 0.5936]])
tensor([[True, True, True, True],
        [True, True, True, True],
        [True, True, True, True]])

```

Extra resources for reproducibility: \* <https://pytorch.org/docs/stable/notes/randomness.html> \* [https://en.wikipedia.org/wiki/Random\\_seed](https://en.wikipedia.org/wiki/Random_seed)

## Running tensors and PyTorch objects on the GPUs (and making faster computations)

GPUs = faster computation on numbers, thanks to CUDA + NVIDIA hardware + PyTorch working behind the scenes to make everything hunky dory (good).

### 0.8.1 1. Getting a GPU

1. Easiest - Use Google Colab for a free GPU (options to upgrade as well)
2. Use your own GPU - takes a little bit of setup and requires the investment of purchasing a GPU, there's lots of options..., see this post for what option to get: <https://timdettmers.com/2020/09/07/which-gpu-for-deep-learning/>
3. Use cloud computing - GCP, AWS, Azure, these services allow you to rent computers on the cloud and access them

For 2, 3 PyTorch + GPU drivers (CUDA) takes a little bit of setting up, to do this, refer to PyTorch setup documentation: <https://pytorch.org/get-started/locally/>

```
[ ]: !nvidia-smi
```

```
Sun Feb 20 00:24:35 2022
```

```
+-----+
| NVIDIA-SMI 460.32.03      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+-----+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan   Temp   Perf   Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                               |                    |            MIG M. |
+-----+-----+-----+-----+-----+-----+
|    0   Tesla P100-PCIE...    Off      | 00000000:00:04:0  Off |              0      |
| N/A    32C    P0      28W / 250W |      0MiB / 16280MiB |      0%      Default |
|                               |                    |              N/A |
+-----+-----+-----+-----+-----+-----+

+-----+
| Processes:
| GPU   GI    CI          PID    Type    Process name                  GPU Memory
|       ID    ID                                   |          Usage
+-----+-----+-----+-----+-----+-----+
| No running processes found
+-----+
```

### 0.8.2 2. Check for GPU access with PyTorch

```
[ ]: # Check for GPU access with PyTorch
import torch
torch.cuda.is_available()
```

```
[ ]: True
```

For PyTorch since it's capable of running compute on the GPU or CPU, it's best practice to setup device agnostic code: <https://pytorch.org/docs/stable/notes/cuda.html#best-practices>

E.g. run on GPU if available, else default to CPU



```
[ ]: # Setup device agnostic code
device = "cuda" if torch.cuda.is_available() else "cpu"
device
```

```
[ ]: 'cuda'
```

```
[ ]: # Count number of devices
torch.cuda.device_count()
```

```
[ ]: 1
```

### 0.9 3. Putting tensors (and models) on the GPU

The reason we want our tensors/models on the GPU is because using a GPU results in faster computations.

```
[ ]: # Create a tensor (default on the CPU)
tensor = torch.tensor([1, 2, 3])

# Tensor not on GPU
print(tensor, tensor.device)
```

```
tensor([1, 2, 3]) cpu
```

```
[ ]: # Move tensor to GPU (if available)
tensor_on_gpu = tensor.to(device)
tensor_on_gpu
```

```
[ ]: tensor([1, 2, 3], device='cuda:0')
```

#### 0.9.1 4. Moving tensors back to the CPU

```
[ ]: # If tensor is on GPU, can't transform it to NumPy
tensor_on_gpu.numpy()
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-8-b7da913938a5> in <module>()
      1 # If tensor is on GPU, can't transform it to NumPy
----> 2 tensor_on_gpu.numpy()

TypeError: can't convert cuda:0 device type tensor to numpy. Use Tensor.cpu() to
↳ copy the tensor to host memory first.
```

```
[ ]: # To fix the GPU tensor with NumPy issue, we can first set it to the CPU
tensor_back_on_cpu = tensor_on_gpu.cpu().numpy()
```

```
tensor_back_on_cpu
```

```
[ ]: array([1, 2, 3])
```

```
[ ]: tensor_on_gpu
```

```
[ ]: tensor([1, 2, 3], device='cuda:0')
```

## 0.10 Exercises & Extra-curriculum

See exercises for this notebook here: [https://www.learnpytorch.io/00\\_pytorch\\_fundamentals/#exercises](https://www.learnpytorch.io/00_pytorch_fundamentals/#exercises)

See the template exercises notebook for this module here: [https://github.com/mrdbourke/pytorch-deep-learning/blob/main/extras/exercises/00\\_pytorch\\_fundamentals\\_exercises.ipynb](https://github.com/mrdbourke/pytorch-deep-learning/blob/main/extras/exercises/00_pytorch_fundamentals_exercises.ipynb)

```
[ ]:
```