

05. Going Modular: Part 1 (cell mode)

This notebook is part 1/2 of section [05. Going Modular](#).

For reference, the two parts are:

1. [05. Going Modular: Part 1 \(cell mode\)](#) - this notebook is run as a traditional Jupyter Notebook/Google Colab notebook and is a condensed version of [notebook 04](#).
2. [05. Going Modular: Part 2 \(script mode\)](#) - this notebook is the same as number 1 but with added functionality to turn each of the major sections into Python scripts, such as, `data_setup.py` and `train.py`.

Why two parts?

Because sometimes the best way to learn something is to see how it *differs* from something else.

If you run each notebook side-by-side you'll see how they differ and that's where the key learnings are.

What is cell mode?

A cell mode notebook is a regular notebook run exactly how we've been running them through the course.

Some cells contain text and others contain code.

What's the difference between this notebook (Part 1) and the script mode notebook (Part 2)?

This notebook, 05. PyTorch Going Modular: Part 1 (cell mode), runs a cleaned up version of the most useful code from section [04. PyTorch Custom Datasets](#).

Running this notebook end-to-end will result in recreating the image classification model we built in notebook 04 (TinyVGG) trained on images of pizza, steak and sushi.

The main difference between this notebook (Part 1) and Part 2 is that each section in Part 2 (script mode) has an extra subsection (e.g. 2.1, 3.1, 4.1) for turning cell code into script code.

Where can you get help?

You can find the book version of this section [05. PyTorch Going Modular on learnpytorch.io](#).

The rest of the materials for this course [are available on GitHub](#).

If you run into trouble, you can ask a question on the course [GitHub Discussions page](#).

And of course, there's the [PyTorch documentation](#) and [PyTorch developer forums](#), a very helpful place for all things PyTorch.

0. Running a notebook in cell mode

As discussed, we're going to be running this notebook normally.

One cell at a time.

The code is from notebook 04, however, it has been condensed down to its core functionality.

1. Get data

We're going to start by downloading the same data we used in [notebook 04](#), the `pizza_steak_sushi` dataset with images of pizza, steak and sushi.

```
import os
import zipfile

from pathlib import Path

import requests

# Setup path to data folder
data_path = Path("data/")
image_path = data_path / "pizza_steak_sushi"

# If the image folder doesn't exist, download it and prepare it...
if image_path.is_dir():
    print(f"{image_path} directory exists.")
else:
    print(f"Did not find {image_path} directory, creating one...")
    image_path.mkdir(parents=True, exist_ok=True)

# Download pizza, steak, sushi data
with open(data_path / "pizza_steak_sushi.zip", "wb") as f:
    request = requests.get("https://github.com/mrdbourke/pytorch-deep-learning/raw/main/data/pizza_steak_sushi.zip")
    print("Downloading pizza, steak, sushi data...")
    f.write(request.content)

# Unzip pizza, steak, sushi data
with zipfile.ZipFile(data_path / "pizza_steak_sushi.zip", "r") as zip_ref:
    print("Unzipping pizza, steak, sushi data...")
    zip_ref.extractall(image_path)

# Remove zip file
os.remove(data_path / "pizza_steak_sushi.zip")

Did not find data/pizza_steak_sushi directory, creating one...
Downloading pizza, steak, sushi data...
Unzipping pizza, steak, sushi data...
```

```
# Setup train and testing paths
train_dir = image_path / "train"
test_dir = image_path / "test"

train_dir, test_dir

(PosixPath('data/pizza_steak_sushi/train'),
 PosixPath('data/pizza_steak_sushi/test'))
```

2. Create Datasets and DataLoaders

Now we'll turn the image dataset into PyTorch `Dataset`'s and `DataLoader`'s.

```
from torchvision import datasets, transforms

# Create simple transform
data_transform = transforms.Compose([
    transforms.Resize((64, 64)),
    transforms.ToTensor(),
])

# Use ImageFolder to create dataset(s)
train_data = datasets.ImageFolder(root=train_dir, # target folder of
                                  # transforms to perform on data (images)
                                  transform=data_transform, #
                                  target_transform=None) # transforms
                                                         # to perform on labels (if necessary)

test_data = datasets.ImageFolder(root=test_dir,
                                  transform=data_transform)

print(f"Train data:\n{train_data}\nTest data:\n{test_data}")

Train data:
Dataset ImageFolder
  Number of datapoints: 225
  Root location: data/pizza_steak_sushi/train
  StandardTransform
Transform: Compose(
  Resize(size=(64, 64), interpolation=bilinear,
max_size=None, antialias=None)
  ToTensor()
)

Test data:
Dataset ImageFolder
  Number of datapoints: 75
  Root location: data/pizza_steak_sushi/test
  StandardTransform
```

```

Transform: Compose(
    Resize(size=(64, 64), interpolation=bilinear,
max_size=None, antialias=None)
    ToTensor()
)

# Get class names as a list
class_names = train_data.classes
class_names

['pizza', 'steak', 'sushi']

# Can also get class names as a dict
class_dict = train_data.class_to_idx
class_dict

{'pizza': 0, 'steak': 1, 'sushi': 2}

# Check the lengths
len(train_data), len(test_data)

(225, 75)

# Turn train and test Datasets into DataLoaders
from torch.utils.data import DataLoader
train_dataloader = DataLoader(dataset=train_data,
                             batch_size=1, # how many samples per
batch?
                             num_workers=1, # how many subprocesses
to use for data loading? (higher = more)
                             shuffle=True) # shuffle the data?

test_dataloader = DataLoader(dataset=test_data,
                             batch_size=1,
                             num_workers=1,
                             shuffle=False) # don't usually need to
shuffle testing data

train_dataloader, test_dataloader

(<torch.utils.data.dataloader.DataLoader at 0x7fed03d16b50>,
 <torch.utils.data.dataloader.DataLoader at 0x7fed03d16590>)

# Check out single image size/shape
img, label = next(iter(train_dataloader))

# Batch size will now be 1, try changing the batch_size parameter
above and see what happens
print(f"Image shape: {img.shape} -> [batch_size, color_channels,
height, width]")
print(f"Label shape: {label.shape}")

```

```
Image shape: torch.Size([1, 3, 64, 64]) -> [batch_size,
color_channels, height, width]
Label shape: torch.Size([1])
```

2.1 Create Datasets and DataLoaders (script mode)

Let's use the Jupyter magic function to create a `.py` file for creating DataLoaders.

We can save a code cell's contents to a file using the Jupyter magic `%%writefile filename` - <https://ipython.readthedocs.io/en/stable/interactive/magics.html#cellmagic-writefile>

```
# Create a directory going_modular scripts
import os
os.makedirs("going_modular")

%%writefile going_modular/data_setup.py
"""
Contains functionality for creating PyTorch DataLoader's for
image classification data.
"""

import os

from torchvision import datasets, transforms
from torch.utils.data import DataLoader

NUM_WORKERS = os.cpu_count()

def create_dataloaders(
    train_dir: str,
    test_dir: str,
    transform: transforms.Compose,
    batch_size: int,
    num_workers: int=NUM_WORKERS
):
    """Creates training and testing DataLoaders.

    Takes in a training directory and testing directory path and turns
    them into
    PyTorch Datasets and then into PyTorch DataLoaders.

    Args:
        train_dir: Path to training directory.
        test_dir: Path to testing directory.
        transform: torchvision transforms to perform on training and
        testing data.
        batch_size: Number of samples per batch in each of the
        DataLoaders.
        num_workers: An integer for number of workers per DataLoader.

    Returns:
```

```

    A tuple of (train_dataloader, test_dataloader, class_names).
    Where class_names is a list of the target classes.
    Example usage:
        train_dataloader, test_dataloader, class_names =
create_data_loaders(train_dir=path/to/train_dir,
                    test_dir=path/to/test_dir,
                    transform=some_transform,
                    batch_size=32,
                    num_workers=4)
"""
# Use ImageFolder to create datasets(s)
train_data = datasets.ImageFolder(train_dir, transform=transform)
test_data = datasets.ImageFolder(test_dir, transform=transform)

# Get class names
class_names = train_data.classes

# Turn images into DataLoaders
train_dataloader = DataLoader(
    train_data,
    batch_size=batch_size,
    shuffle=True,
    num_workers=num_workers,
    pin_memory=True # for more on pin memory, see the PyTorch docs:
https://pytorch.org/docs/stable/data.html
)

test_dataloader = DataLoader(
    test_data,
    batch_size=batch_size,
    shuffle=False,
    num_workers=num_workers,
    pin_memory=True
)

return train_dataloader, test_dataloader, class_names

```

Overwriting going_modular/data_setup.py

```

from going_modular import data_setup

train_dataloader, test_dataloader, class_names =
data_setup.create_data_loaders(train_dir=train_dir,

test_dir=test_dir,

transform=data_transform,

batch_size=32)
train_dataloader, test_dataloader, class_names

```

```
(<torch.utils.data.dataloader.DataLoader at 0x7fec47a09490>,
<torch.utils.data.dataloader.DataLoader at 0x7fec47a09610>,
['pizza', 'steak', 'sushi'])
```

3. Making a model (TinyVGG)

We're going to use the same model we used in notebook 04: TinyVGG from the CNN Explainer website.

The only change here from notebook 04 is that a docstring has been added using [Google's Style Guide for Python](#).

```
import torch

from torch import nn

class TinyVGG(nn.Module):
    """Creates the TinyVGG architecture.

    Replicates the TinyVGG architecture from the CNN explainer website
    in PyTorch.
    See the original architecture here: https://poloclub.github.io/cnn-explainer/

    Args:
        input_shape: An integer indicating number of input channels.
        hidden_units: An integer indicating number of hidden units between
        layers.
        output_shape: An integer indicating number of output units.
    """
    def __init__(self, input_shape: int, hidden_units: int,
output_shape: int) -> None:
        super().__init__()
        self.conv_block_1 = nn.Sequential(
            nn.Conv2d(in_channels=input_shape,
                      out_channels=hidden_units,
                      kernel_size=3, # how big is the square that's
going over the image?
                      stride=1, # default
                      padding=0), # options = "valid" (no padding) or
"same" (output has same shape as input) or int for specific number
            nn.ReLU(),
            nn.Conv2d(in_channels=hidden_units,
                      out_channels=hidden_units,
                      kernel_size=3,
                      stride=1,
                      padding=0),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2,
```

```

        stride=2) # default stride value is same as
kernel_size
    )
    self.conv_block_2 = nn.Sequential(
        nn.Conv2d(hidden_units, hidden_units, kernel_size=3,
padding=0),
        nn.ReLU(),
        nn.Conv2d(hidden_units, hidden_units, kernel_size=3,
padding=0),
        nn.ReLU(),
        nn.MaxPool2d(2)
    )
    self.classifier = nn.Sequential(
        nn.Flatten(),
        # Where did this in_features shape come from?
        # It's because each layer of our network compresses and
changes the shape of our inputs data.
        nn.Linear(in_features=hidden_units*13*13,
                    out_features=output_shape)
    )

    def forward(self, x: torch.Tensor):
        x = self.conv_block_1(x)
        x = self.conv_block_2(x)
        x = self.classifier(x)
        return x
        # return self.classifier(self.block_2(self.block_1(x))) # <-
leverage the benefits of operator fusion

import torch

device = "cuda" if torch.cuda.is_available() else "cpu"

# Instantiate an instance of the model
torch.manual_seed(42)
model_0 = TinyVGG(input_shape=3, # number of color channels (3 for
RGB)
                    hidden_units=10,
                    output_shape=len(train_data.classes)).to(device)

model_0
TinyVGG(
  (conv_block_1): Sequential(
    (0): Conv2d(3, 10, kernel_size=(3, 3), stride=(1, 1))
    (1): ReLU()
    (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )

```



```

    (conv_block_2): Sequential(
      (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
      (1): ReLU()
      (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
      (3): ReLU()
      (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (classifier): Sequential(
      (0): Flatten(start_dim=1, end_dim=-1)
      (1): Linear(in_features=1690, out_features=3, bias=True)
    )
  )
)

```

To test our model let's do a single forward pass (pass a sample batch from the training set through our model).

```

# 1. Get a batch of images and labels from the DataLoader
img_batch, label_batch = next(iter(train_dataloader))

# 2. Get a single image from the batch and unsqueeze the image so its
shape fits the model
img_single, label_single = img_batch[0].unsqueeze(dim=0),
label_batch[0]
print(f"Single image shape: {img_single.shape}\n")

# 3. Perform a forward pass on a single image
model_0.eval()
with torch.inference_mode():
    pred = model_0(img_single.to(device))

# 4. Print out what's happening and convert model logits -> pred probs
-> pred label
print(f"Output logits:\n{pred}\n")
print(f"Output prediction probabilities:\n{torch.softmax(pred,
dim=1)}\n")
print(f"Output prediction label:\n{torch.argmax(torch.softmax(pred,
dim=1), dim=1)}\n")
print(f"Actual label:\n{label_single}")

Single image shape: torch.Size([1, 3, 64, 64])

Output logits:
tensor([[ 0.0208, -0.0019,  0.0095]], device='cuda:0')

Output prediction probabilities:
tensor([[0.3371, 0.3295, 0.3333]], device='cuda:0')

Output prediction label:
tensor([0], device='cuda:0')

```

Actual label:
0

3.1. Making a model (TinyVGG) with a script (model_builder.py)

Let's turn our model building code into a Python script we can import.

```
%%writefile going_modular/model_builder.py
"""
Contains PyTorch model code to instantiate a TinyVGG model from the
CNN Explainer website.
"""
import torch

from torch import nn

class TinyVGG(nn.Module):
    """Creates the TinyVGG architecture.

    Replicates the TinyVGG architecture from the CNN explainer website
in PyTorch.
    See the original architecture here: https://poloclub.github.io/cnn-explainer/

    Args:
        input_shape: An integer indicating number of input channels.
        hidden_units: An integer indicating number of hidden units between
        layers.
        output_shape: An integer indicating number of output units.
    """
    def __init__(self, input_shape: int, hidden_units: int,
output_shape: int) -> None:
        super().__init__()
        self.conv_block_1 = nn.Sequential(
            nn.Conv2d(in_channels=input_shape,
                    out_channels=hidden_units,
                    kernel_size=3, # how big is the square that's
going over the image?
                    stride=1, # default
padding=0), # options = "valid" (no padding) or
"same" (output has same shape as input) or int for specific number
            nn.ReLU(),
            nn.Conv2d(in_channels=hidden_units,
                    out_channels=hidden_units,
                    kernel_size=3,
                    stride=1,
                    padding=0),
            nn.ReLU(),
```

```

        nn.MaxPool2d(kernel_size=2,
                      stride=2) # default stride value is same as
kernel_size
    )
    self.conv_block_2 = nn.Sequential(
        nn.Conv2d(hidden_units, hidden_units, kernel_size=3,
padding=0),
        nn.ReLU(),
        nn.Conv2d(hidden_units, hidden_units, kernel_size=3,
padding=0),
        nn.ReLU(),
        nn.MaxPool2d(2)
    )
    self.classifier = nn.Sequential(
        nn.Flatten(),
        # Where did this in_features shape come from?
        # It's because each layer of our network compresses and
changes the shape of our inputs data.
        nn.Linear(in_features=hidden_units*13*13,
                    out_features=output_shape)
    )

    def forward(self, x: torch.Tensor):
        x = self.conv_block_1(x)
        x = self.conv_block_2(x)
        x = self.classifier(x)
        return x
        # return self.classifier(self.block_2(self.block_1(x))) # <-
leverage the benefits of operator fusion

```

Overwriting going_modular/model_builder.py

```

from going_modular import model_builder

import torch

from going_modular import model_builder

device = "cuda" if torch.cuda.is_available() else "cpu"

# Instantiate a model from the model_builder.py script
torch.manual_seed(42)
model_1 = model_builder.TinyVGG(input_shape=3,
                                hidden_units=10,

output_shape=len(class_names)).to(device)
model_1

TinyVGG(
  (conv_block_1): Sequential(
    (0): Conv2d(3, 10, kernel_size=(3, 3), stride=(1, 1))

```

```

        (1): ReLU()
        (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
        (3): ReLU()
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (conv_block_2): Sequential(
      (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
      (1): ReLU()
      (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
      (3): ReLU()
      (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (classifier): Sequential(
      (0): Flatten(start_dim=1, end_dim=-1)
      (1): Linear(in_features=1690, out_features=3, bias=True)
    )
  )

```

1. Get a batch of images and labels from the DataLoader

2. Get a single image from the batch and unsqueeze the image so its shape fits the model

label_batch[0]
print(f"Single image shape: {img_single.shape}\n")

3. Perform a forward pass on a single image
model_1.eval()
with torch.inference_mode():
 pred = model_1(img_single.to(device))

4. Print out what's happening and convert model logits -> pred probs -> pred label
print(f"Output logits:\n{pred}\n")
print(f"Output prediction probabilities:\n{torch.softmax(pred,
dim=1)}\n")
print(f"Output prediction label:\n{torch.argmax(torch.softmax(pred,
dim=1), dim=1)}\n")
print(f"Actual label:\n{label_single}")

Single image shape: torch.Size([1, 3, 64, 64])

Output logits:
tensor([[0.0208, -0.0019, 0.0095]], device='cuda:0')

Output prediction probabilities:
tensor([[0.3371, 0.3295, 0.3333]], device='cuda:0')

```
Output prediction label:  
tensor([0], device='cuda:0')
```

```
Actual label:  
0
```

4. Creating `train_step()` and `test_step()` functions and `train()` to combine them

Rather than writing them again, we can reuse the `train_step()` and `test_step()` functions from [notebook 04](#).

The same goes for the `train()` function we created.

The only difference here is that these functions have had docstrings added to them in [Google's Python Functions and Methods Style Guide](#).

Let's start by making `train_step()`.

```
from typing import Tuple  
  
def train_step(model: torch.nn.Module,  
               dataloader: torch.utils.data.DataLoader,  
               loss_fn: torch.nn.Module,  
               optimizer: torch.optim.Optimizer,  
               device: torch.device) -> Tuple[float, float]:  
    """Trains a PyTorch model for a single epoch.  
  
    Turns a target PyTorch model to training mode and then  
    runs through all of the required training steps (forward  
    pass, loss calculation, optimizer step).  
  
    Args:  
        model: A PyTorch model to be trained.  
        dataloader: A DataLoader instance for the model to be trained on.  
        loss_fn: A PyTorch loss function to minimize.  
        optimizer: A PyTorch optimizer to help minimize the loss function.  
        device: A target device to compute on (e.g. "cuda" or "cpu").  
  
    Returns:  
        A tuple of training loss and training accuracy metrics.  
        In the form (train_loss, train_accuracy). For example:  
  
        (0.1112, 0.8743)  
    """  
    # Put model in train mode  
    model.train()
```

```

# Setup train loss and train accuracy values
train_loss, train_acc = 0, 0

# Loop through data loader data batches
for batch, (X, y) in enumerate(dataloader):
    # Send data to target device
    X, y = X.to(device), y.to(device)

    # 1. Forward pass
    y_pred = model(X)

    # 2. Calculate and accumulate loss
    loss = loss_fn(y_pred, y)
    train_loss += loss.item()

    # 3. Optimizer zero grad
    optimizer.zero_grad()

    # 4. Loss backward
    loss.backward()

    # 5. Optimizer step
    optimizer.step()

    # Calculate and accumulate accuracy metric across all batches
    y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim=1)
    train_acc += (y_pred_class == y).sum().item()/len(y_pred)

# Adjust metrics to get average loss and accuracy per batch
train_loss = train_loss / len(dataloader)
train_acc = train_acc / len(dataloader)
return train_loss, train_acc

```

Now we'll do `test_step()`.

```

def test_step(model: torch.nn.Module,
              dataloader: torch.utils.data.DataLoader,
              loss_fn: torch.nn.Module,
              device: torch.device) -> Tuple[float, float]:
    """Tests a PyTorch model for a single epoch.

    Turns a target PyTorch model to "eval" mode and then performs
    a forward pass on a testing dataset.

    Args:
        model: A PyTorch model to be tested.
        dataloader: A DataLoader instance for the model to be tested on.
        loss_fn: A PyTorch loss function to calculate loss on the test
        data.
        device: A target device to compute on (e.g. "cuda" or "cpu").
    """

```

Returns:

A tuple of testing loss and testing accuracy metrics.
In the form (test_loss, test_accuracy). For example:

(0.0223, 0.8985)

"""

Put model in eval mode

model.eval()

Setup test loss and test accuracy values

test_loss, test_acc = 0, 0

Turn on inference context manager

with torch.inference_mode():

Loop through DataLoader batches

for batch, (X, y) in enumerate(dataloader):

Send data to target device

X, y = X.to(device), y.to(device)

1. Forward pass

test_pred_logits = model(X)

2. Calculate and accumulate loss

loss = loss_fn(test_pred_logits, y)

test_loss += loss.item()

Calculate and accumulate accuracy

test_pred_labels = test_pred_logits.argmax(dim=1)

test_acc += ((test_pred_labels ==

y).sum()).item()/len(test_pred_labels))

Adjust metrics to get average loss and accuracy per batch

test_loss = test_loss / len(dataloader)

test_acc = test_acc / len(dataloader)

return test_loss, test_acc

And we'll combine `train_step()` and `test_step()` into `train()`.

```
from typing import Dict, List
```

```
from tqdm.auto import tqdm
```

```
def train(model: torch.nn.Module,
          train_dataloader: torch.utils.data.DataLoader,
          test_dataloader: torch.utils.data.DataLoader,
          optimizer: torch.optim.Optimizer,
          loss_fn: torch.nn.Module,
          epochs: int,
          device: torch.device) -> Dict[str, List[float]]:
```

```
data_loader=test_data_loader,
```



```

        loss_fn=loss_fn,
        device=device)

    # Print out what's happening
    print(
        f"Epoch: {epoch+1} | "
        f"train_loss: {train_loss:.4f} | "
        f"train_acc: {train_acc:.4f} | "
        f"test_loss: {test_loss:.4f} | "
        f"test_acc: {test_acc:.4f}"
    )

    # Update results dictionary
    results["train_loss"].append(train_loss)
    results["train_acc"].append(train_acc)
    results["test_loss"].append(test_loss)
    results["test_acc"].append(test_acc)

    # Return the filled results at the end of the epochs
    return results

```

4.1 Turn training functions into a script (engine.py)

```

%%writefile going_modular/engine.py
"""
Contains functions for training and testing a PyTorch model.
"""
from typing import Dict, List, Tuple

import torch

from tqdm.auto import tqdm

def train_step(model: torch.nn.Module,
               dataloader: torch.utils.data.DataLoader,
               loss_fn: torch.nn.Module,
               optimizer: torch.optim.Optimizer,
               device: torch.device) -> Tuple[float, float]:
    """Trains a PyTorch model for a single epoch.

    Turns a target PyTorch model to training mode and then
    runs through all of the required training steps (forward
    pass, loss calculation, optimizer step).

    Args:
        model: A PyTorch model to be trained.
        dataloader: A DataLoader instance for the model to be trained on.
        loss_fn: A PyTorch loss function to minimize.
        optimizer: A PyTorch optimizer to help minimize the loss function.
        device: A target device to compute on (e.g. "cuda" or "cpu").
    """

```

Returns:

A tuple of training loss and training accuracy metrics.
In the form (train_loss, train_accuracy). For example:

```
(0.1112, 0.8743)
"""
# Put model in train mode
model.train()

# Setup train loss and train accuracy values
train_loss, train_acc = 0, 0

# Loop through data loader data batches
for batch, (X, y) in enumerate(dataloader):
    # Send data to target device
    X, y = X.to(device), y.to(device)

    # 1. Forward pass
    y_pred = model(X)

    # 2. Calculate and accumulate loss
    loss = loss_fn(y_pred, y)
    train_loss += loss.item()

    # 3. Optimizer zero grad
    optimizer.zero_grad()

    # 4. Loss backward
    loss.backward()

    # 5. Optimizer step
    optimizer.step()

    # Calculate and accumulate accuracy metric across all batches
    y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim=1)
    train_acc += (y_pred_class == y).sum().item()/len(y_pred)

# Adjust metrics to get average loss and accuracy per batch
train_loss = train_loss / len(dataloader)
train_acc = train_acc / len(dataloader)
return train_loss, train_acc

def test_step(model: torch.nn.Module,
              dataloader: torch.utils.data.DataLoader,
              loss_fn: torch.nn.Module,
              device: torch.device) -> Tuple[float, float]:
    """Tests a PyTorch model for a single epoch.

    Turns a target PyTorch model to "eval" mode and then performs
```

a forward pass on a testing dataset.

Args:

model: A PyTorch model to be tested.

dataloader: A DataLoader instance for the model to be tested on.

loss_fn: A PyTorch loss function to calculate loss on the test data.

device: A target device to compute on (e.g. "cuda" or "cpu").

Returns:

A tuple of testing loss and testing accuracy metrics.

In the form (test_loss, test_accuracy). For example:

(0.0223, 0.8985)

"""

Put model in eval mode

model.eval()

Setup test loss and test accuracy values

test_loss, test_acc = 0, 0

Turn on inference context manager

with torch.inference_mode():

Loop through DataLoader batches

for batch, (X, y) in enumerate(dataloader):

Send data to target device

X, y = X.to(device), y.to(device)

1. Forward pass

test_pred_logits = model(X)

2. Calculate and accumulate loss

loss = loss_fn(test_pred_logits, y)

test_loss += loss.item()

Calculate and accumulate accuracy

test_pred_labels = test_pred_logits.argmax(dim=1)

test_acc += ((test_pred_labels ==
y).sum().item()/len(test_pred_labels))

Adjust metrics to get average loss and accuracy per batch

test_loss = test_loss / len(dataloader)

test_acc = test_acc / len(dataloader)

return test_loss, test_acc

```
def train(model: torch.nn.Module,  
          train_dataloader: torch.utils.data.DataLoader,  
          test_dataloader: torch.utils.data.DataLoader,  
          optimizer: torch.optim.Optimizer,
```



```

                                device=device)
test_loss, test_acc = test_step(model=model,
                                dataloader=test_dataloader,
                                loss_fn=loss_fn,
                                device=device)

# Print out what's happening
print(
    f"Epoch: {epoch+1} | "
    f"train_loss: {train_loss:.4f} | "
    f"train_acc: {train_acc:.4f} | "
    f"test_loss: {test_loss:.4f} | "
    f"test_acc: {test_acc:.4f}"
)

# Update results dictionary
results["train_loss"].append(train_loss)
results["train_acc"].append(train_acc)
results["test_loss"].append(test_loss)
results["test_acc"].append(test_acc)

# Return the filled results at the end of the epochs
return results

```

Writing `going_modular/engine.py`

```

from going_modular import engine

# engine.train()

```

5. Creating a function to save the model

Let's setup a function to save our model to a directory.

```

from pathlib import Path

def save_model(model: torch.nn.Module,
               target_dir: str,
               model_name: str):
    """Saves a PyTorch model to a target directory.

    Args:
        model: A target PyTorch model to save.
        target_dir: A directory for saving the model to.
        model_name: A filename for the saved model. Should include
            either ".pth" or ".pt" as the file extension.

    Example usage:
        save_model(model=model_0,
                   target_dir="models",

```

```

        model_name="05_going_modular_tingvgg_model.pth")
    """
    # Create target directory
    target_dir_path = Path(target_dir)
    target_dir_path.mkdir(parents=True,
                          exist_ok=True)

    # Create model save path
    assert model_name.endswith(".pth") or model_name.endswith(".pt"),
    "model_name should end with '.pt' or '.pth'"
    model_save_path = target_dir_path / model_name

    # Save the model state_dict()
    print(f"[INFO] Saving model to: {model_save_path}")
    torch.save(obj=model.state_dict(),
               f=model_save_path)

```

5.1 Create a file called `utils.py` with utility functions

"utils" in Python is generally reserved for various utility functions.

Right now we only have one utility function (`save_model()`) but... as our code grows we'll likely have more...

```

%%writefile going_modular/utils.py
"""
File containing various utility functions for PyTorch model training.
"""
import torch

from pathlib import Path

def save_model(model: torch.nn.Module,
               target_dir: str,
               model_name: str):
    """Saves a PyTorch model to a target directory.

    Args:
        model: A target PyTorch model to save.
        target_dir: A directory for saving the model to.
        model_name: A filename for the saved model. Should include
            either ".pth" or ".pt" as the file extension.

    Example usage:
        save_model(model=model_0,
                  target_dir="models",
                  model_name="05_going_modular_tingvgg_model.pth")
    """
    # Create target directory
    target_dir_path = Path(target_dir)

```

```

target_dir_path.mkdir(parents=True,
                      exist_ok=True)

# Create model save path
assert model_name.endswith(".pth") or model_name.endswith(".pt"),
"model_name should end with '.pt' or '.pth'"
model_save_path = target_dir_path / model_name

# Save the model state_dict()
print(f"[INFO] Saving model to: {model_save_path}")
torch.save(obj=model.state_dict(),
           f=model_save_path)

```

Overwriting going_modular/utils.py

6. Train, evaluate and save the model

Let's leverage the functions we've got above to train, test and save a model to file.

```

# Set random seeds
torch.manual_seed(42)
torch.cuda.manual_seed(42)

# Set number of epochs
NUM_EPOCHS = 5

# Recreate an instance of TinyVGG
model_0 = TinyVGG(input_shape=3, # number of color channels (3 for
RGB)
                  hidden_units=10,
                  output_shape=len(train_data.classes)).to(device)

# Setup loss function and optimizer
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(params=model_0.parameters(), lr=0.001)

# Start the timer
from timeit import default_timer as timer
start_time = timer()

# Train model_0
model_0_results = train(model=model_0,
                        train_dataloader=train_dataloader,
                        test_dataloader=test_dataloader,
                        optimizer=optimizer,
                        loss_fn=loss_fn,
                        epochs=NUM_EPOCHS,
                        device=device)

# End the timer and print out how long it took

```

```

end_time = timer()
print(f"[INFO] Total training time: {end_time-start_time:.3f}
seconds")

# Save the model
save_model(model=model_0,
            target_dir="models",
            model_name="05_going_modular_cell_mode_tinyvgg_model.pth")

{"model_id": "ac11b3a31c254a448864f2bfce323e57", "version_major": 2, "version_minor": 0}

Epoch: 1 | train_loss: 1.0962 | train_acc: 0.3778 | test_loss: 1.0684
| test_acc: 0.4133
Epoch: 2 | train_loss: 1.0254 | train_acc: 0.5111 | test_loss: 1.0144
| test_acc: 0.4400
Epoch: 3 | train_loss: 0.9479 | train_acc: 0.5333 | test_loss: 0.9846
| test_acc: 0.4800
Epoch: 4 | train_loss: 0.9026 | train_acc: 0.5778 | test_loss: 0.9910
| test_acc: 0.4400
Epoch: 5 | train_loss: 0.8758 | train_acc: 0.6178 | test_loss: 1.0112
| test_acc: 0.5333
[INFO] Total training time: 13.737 seconds
[INFO] Saving model to:
models/05_going_modular_cell_mode_tinyvgg_model.pth

```

6.1 Train, evaluate and save the model (script mode) -> train.py

Let's create a file called `train.py` to leverage all of our other code scripts to train a PyTorch model.

Essentially we want to replicate the functionality of notebook 04 in one line...

```

%%writefile going_modular/train.py
"""
Trains a PyTorch image classification model using device-agnostic
code.
"""

import os
import torch

from torchvision import transforms
from timeit import default_timer as timer

import data_setup, engine, model_builder, utils

# Setup hyperparameters
NUM_EPOCHS = 5
BATCH_SIZE = 32
HIDDEN_UNITS = 10

```



```

LEARNING_RATE = 0.001

# Setup directories
train_dir = "data/pizza_steak_sushi/train"
test_dir = "data/pizza_steak_sushi/test"

# Setup device agnostic code
device = "cuda" if torch.cuda.is_available() else "cpu"

# Create transforms
data_transform = transforms.Compose([
    transforms.Resize((64, 64)),
    transforms.ToTensor()
])

# Create DataLoader's and get class_names
train_dataloader, test_dataloader, class_names =
data_setup.create_data_loaders(train_dir=train_dir,

test_dir=test_dir,

transform=data_transform,

batch_size=BATCH_SIZE)

# Create model
model = model_builder.TinyVGG(input_shape=3,
                               hidden_units=HIDDEN_UNITS,

output_shape=len(class_names)).to(device)

# Setup loss and optimizer
loss_fn = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(),
                               lr=LEARNING_RATE)

# Start the timer
start_time = timer()

# Start training with help from engine.py
engine.train(model=model,
             train_dataloader=train_dataloader,
             test_dataloader=test_dataloader,
             loss_fn=loss_fn,
             optimizer=optimizer,
             epochs=NUM_EPOCHS,
             device=device)

# End the timer and print out how long it took

```

```

end_time = timer()
print(f"[INFO] Total training time: {end_time-start_time:.3f}
seconds")

# Save the model to file
utils.save_model(model=model,
                  target_dir="models",

model_name="05_going_modular_script_mode_tinyvgg_model.pth")

Overwriting going_modular/train.py

!python going_modular/train.py

 0% 0/5 [00:00<?, ?it/s]Epoch: 1 | train_loss: 1.1019 | train_acc:
0.3203 | test_loss: 1.1089 | test_acc: 0.1979
 20% 1/5 [00:01<00:06, 1.51s/it]Epoch: 2 | train_loss: 1.1055 |
train_acc: 0.2930 | test_loss: 1.1385 | test_acc: 0.1979
 40% 2/5 [00:02<00:04, 1.47s/it]Epoch: 3 | train_loss: 1.1128 |
train_acc: 0.2930 | test_loss: 1.1272 | test_acc: 0.1979
 60% 3/5 [00:04<00:02, 1.45s/it]Epoch: 4 | train_loss: 1.0973 |
train_acc: 0.3906 | test_loss: 1.0982 | test_acc: 0.3021
 80% 4/5 [00:05<00:01, 1.45s/it]Epoch: 5 | train_loss: 1.0959 |
train_acc: 0.4141 | test_loss: 1.1004 | test_acc: 0.3125
100% 5/5 [00:07<00:00, 1.43s/it]
[INFO] Total training time: 7.133 seconds
[INFO] Saving model to:
models/05_going_modular_script_mode_tinyvgg_model.pth

```

We finish with a saved image classification model at
models/05_going_modular_cell_mode_tinyvgg_model.pth.

See exercises and extra-curriculum in section 05 of the Learn PyTorch book -
https://www.learnpytorch.io/05_pytorch_going_modular/#exercises