

07 PyTorch Experiment Tracking

Machine learning is very experimental.

In order to figure out which experiments are worth pursuing, that's where **experiment tracking** comes in, it helps you to figure out what doesn't work so you can figure out what **does** work.

In this notebook, we're going to see an example of programmatically tracking experiments.

Resources:

- Book version of notebook: https://www.learnpytorch.io/07_pytorch_experiment_tracking/
- Ask a question: <https://github.com/mrdbourke/pytorch-deep-learning/discussions>
- Extra-curriculum: <https://madewithml.com/courses/mlops/experiment-tracking/>

```
import torch
import torchvision

print(torch.__version__)
print(torchvision.__version__)

1.13.0.dev20220627+cu113
0.14.0.dev20220627+cu113

# For this notebook to run with updated APIs, we need torch 1.12+ and
# torchvision 0.13+
try:
    import torch
    import torchvision
    assert int(torch.__version__.split(".")[1]) >= 12, "torch version
    should be 1.12+"
    assert int(torchvision.__version__.split(".")[1]) >= 13,
    "torchvision version should be 0.13+"
    print(f"torch version: {torch.__version__}")
    print(f"torchvision version: {torchvision.__version__}")
except:
    print(f"[INFO] torch/torchvision versions not as required,
    installing nightly versions.")
    !pip3 install -U --pre torch torchvision torchaudio --extra-index-
    url https://download.pytorch.org/whl/nightly/cu113
    import torch
    import torchvision
    print(f"torch version: {torch.__version__}")
    print(f"torchvision version: {torchvision.__version__}")

torch version: 1.13.0.dev20220627+cu113
torchvision version: 0.14.0.dev20220627+cu113
```

```

# Continue with regular imports
import matplotlib.pyplot as plt
import torch
import torchvision

from torch import nn
from torchvision import transforms

# Try to get torchinfo, install it if it doesn't work
try:
    from torchinfo import summary
except:
    print("[INFO] Couldn't find torchinfo... installing it.")
    !pip install -q torchinfo
    from torchinfo import summary

# Try to import the going_modular directory, download it from GitHub
if it doesn't work
try:
    from going_modular.going_modular import data_setup, engine
except:
    # Get the going_modular scripts
    print("[INFO] Couldn't find going_modular scripts... downloading
them from GitHub.")
    !git clone https://github.com/mrdbourke/pytorch-deep-learning
    !mv pytorch-deep-learning/going_modular .
    !rm -rf pytorch-deep-learning
    from going_modular.going_modular import data_setup, engine

# Setup device agnostic code
device = "cuda" if torch.cuda.is_available() else "cpu"
device

{"type": "string"}

# Set seeds
def set_seeds(seed: int=42):
    """Sets random seeds for torch operations.

    Args:
        seed (int, optional): Random seed to set. Defaults to 42.
    """
    # Set the seed for general torch operations
    torch.manual_seed(seed)
    # Set the seed for CUDA torch operations (ones that happen on the
GPU)
    torch.cuda.manual_seed(seed)

set_seeds()

```

1. Get data

Want to get pizza, steak, sushi images.

So we can run experiments building FoodVision Mini and see which model performs best.

```
import os
import zipfile

from pathlib import Path

import requests

# example source:
https://github.com/mrdbourke/pytorch-deep-learning/raw/main/data/pizza\_steak\_sushi.zip

def download_data(source: str,
                  destination: str,
                  remove_source: bool = True) -> Path:
    """Downloads a zipped dataset from source and unzips to
    destination."""
    # Setup path to data folder
    data_path = Path("data/")
    image_path = data_path / destination

    # If the image folder doesn't exist, create it
    if image_path.is_dir():
        print(f"[INFO] {image_path} directory already exists, skipping
        download.")
    else:
        print(f"[INFO] Did not find {image_path} directory, creating
        one...")
        image_path.mkdir(parents=True, exist_ok=True)

    # Download the target data
    target_file = Path(source).name
    with open(data_path / target_file, "wb") as f:
        request = requests.get(source)
        print(f"[INFO] Downloading {target_file} from {source}...")
        f.write(request.content)

    # Unzip target file
    with zipfile.ZipFile(data_path / target_file, "r") as zip_ref:
        print(f"[INFO] Unzipping {target_file} data...")
        zip_ref.extractall(image_path)

    # Remove .zip file if needed
    if remove_source:
        os.remove(data_path / target_file)
```

```

    return image_path

image_path =
download_data(source="https://github.com/mrdbourke/pytorch-deep-
learning/raw/main/data/pizza_steak_sushi.zip",
              destination="pizza_steak_sushi")
image_path

[INFO] data/pizza_steak_sushi directory already exists, skipping
download.

PosixPath('data/pizza_steak_sushi')

```

2. Create Datasets and DataLoaders

2.1 Create DataLoaders with manual transforms

The goal with transforms is to ensure your custom data is formatted in a reproducible way as well as a way that will suit pretrained models.

```

# Setup directories
train_dir = image_path / "train"
test_dir = image_path / "test"

train_dir, test_dir

(PosixPath('data/pizza_steak_sushi/train'),
 PosixPath('data/pizza_steak_sushi/test'))

# Setup ImageNet normalization levels
# See here: https://pytorch.org/vision/0.12/models.html
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])

# Create transform pipeline manually
from torchvision import transforms
manual_transforms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    normalize
])
print(f"Manually created transforms: {manual_transforms}")

# Create DataLoaders
from going_modular.going_modular import data_setup
train_dataloader, test_dataloader, class_names =
data_setup.create_data_loaders(train_dir=train_dir,
test_dir=test_dir,

```

```

transform=manual_transforms,

batch_size=32)
train_dataloader, test_dataloader, class_names

Manually created transforms: Compose(
  Resize(size=(224, 224), interpolation=bilinear, max_size=None,
  antialias=None)
  ToTensor()
  Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
)

(<torch.utils.data.dataloader.DataLoader at 0x7fe2dad1590>,
 <torch.utils.data.dataloader.DataLoader at 0x7fe2dad16d0>,
 ['pizza', 'steak', 'sushi'])

```

2.2 Create DataLoaders using automatically created transforms

The same principle applies for automatic transforms: we want our custom data in the same format as a pretrained model was trained on.

```

# Setup dirs
train_dir = image_path / "train"
test_dir = image_path / "test"

# Setup pretrained weights (plenty of these weights available in
torchvision.models v0.13+)
import torchvision
weights = torchvision.models.EfficientNet_B0_Weights.DEFAULT #
"DEFAULT" = best available

# Get transforms from weights (these are the transforms used to train
a particular or obtain a particular set of weights)
automatic_transforms = weights.transforms()
print(f"Automatically created transforms: {automatic_transforms}")

# Create DataLoaders
train_dataloader, test_dataloader, class_names =
data_setup.create_data_loaders(train_dir=train_dir,

test_dir=test_dir,

transform=automatic_transforms,

batch_size=32)
train_dataloader, test_dataloader, class_names

Automatically created transforms: ImageClassification(
  crop_size=[224]
)

```

```

        resize_size=[256]
        mean=[0.485, 0.456, 0.406]
        std=[0.229, 0.224, 0.225]
        interpolation=InterpolationMode.BICUBIC
    )

(<torch.utils.data.dataloader.DataLoader at 0x7fe2dae4ae90>,
 <torch.utils.data.dataloader.DataLoader at 0x7fe2dae07890>,
 ['pizza', 'steak', 'sushi'])

```

3. Getting a pretrained model, freeze the base layers and change the classifier head

```

# Note: This is how a pretrained model would be created prior to
torchvision v0.13
# model =
torchvision.models.efficientnet_b0(pretrained=True).to(device) # OLD

# Download the pretrained weights for EfficientNet_B0
weights = torchvision.models.EfficientNet_B0_Weights.DEFAULT #
"DEFAULT" = best available weights

# Setup the model with the pretrained weights and send it to the
target device
model = torchvision.models.efficientnet_b0(weights=weights).to(device)
# model

Downloading:
"https://download.pytorch.org/models/efficientnet_b0_rwightman-
3dd342df.pth" to
/root/.cache/torch/hub/checkpoints/efficientnet_b0_rwightman-
3dd342df.pth

{"model_id": "98cad9ec7c6b48dc91f3dfa151019292", "version_major": 2, "vers
ion_minor": 0}

# Freeze all base layers by setting their requires_grad attribute to
False
for param in model.features.parameters():
    # print(param)
    param.requires_grad = False

# Adjust the classifier head
set_seeds()
model.classifier = nn.Sequential(
    nn.Dropout(p=0.2, inplace=True),
    nn.Linear(in_features=1280,
out_features=len(class_names)).to(device)

```

```
from torchinfo import summary
```

```
summary(model,
        input_size=(32, 3, 224, 224),
        verbose=0,
        col_names=["input_size", "output_size", "num_params",
"trainable"],
        col_width=20,
        row_settings=["var_names"])
```

```
=====
Layer (type (var_name))      Input
Shape      Output Shape      Param #      Trainable
=====
EfficientNet (EfficientNet)   [32, 3,
224, 224]      [32, 3]      --      Partial
└─Sequential (features)      [32, 3,
224, 224]      [32, 1280, 7, 7]      --      False
|   └─Conv2dNormActivation (0) [32, 3,
224, 224]      [32, 32, 112, 112]      --      False
|   |   └─Conv2d (0)          [32, 3,
224, 224]      [32, 32, 112, 112]      (864)      False
|   |   |   └─BatchNorm2d (1) [32, 32,
112, 112]      [32, 32, 112, 112]      (64)      False
|   |   |   |   └─SiLU (2)    [32, 32,
112, 112]      [32, 32, 112, 112]      --      --
|   |   |   |   └─Sequential (1) [32, 32,
112, 112]      [32, 16, 112, 112]      --      False
|   |   |   |   |   └─MBConv (0) [32, 32,
112, 112]      [32, 16, 112, 112]      (1,448)      False
|   |   |   |   |   |   └─Sequential (2) [32, 16,
112, 112]      [32, 24, 56, 56]      --      False
|   |   |   |   |   |   |   └─MBConv (0) [32, 16,
112, 112]      [32, 24, 56, 56]      (6,004)      False
|   |   |   |   |   |   |   |   └─MBConv (1) [32, 24,
56, 56]      [32, 24, 56, 56]      (10,710)      False
|   |   |   |   |   |   |   |   |   └─Sequential (3) [32, 24,
56, 56]      [32, 40, 28, 28]      --      False
|   |   |   |   |   |   |   |   |   |   └─MBConv (0) [32, 24,
56, 56]      [32, 40, 28, 28]      (15,350)      False
|   |   |   |   |   |   |   |   |   |   |   └─MBConv (1) [32, 40,
28, 28]      [32, 40, 28, 28]      (31,290)      False
|   |   |   |   |   |   |   |   |   |   |   |   └─Sequential (4) [32, 40,
28, 28]      [32, 80, 14, 14]      --      False
|   |   |   |   |   |   |   |   |   |   |   |   |   └─MBConv (0) [32, 40,
28, 28]      [32, 80, 14, 14]      (37,130)      False
|   |   |   |   |   |   |   |   |   |   |   |   |   |   └─MBConv (1) [32, 80,
14, 14]      [32, 80, 14, 14]      (102,900)      False
```

		└─MBConv (2)			[32, 80,
14, 14]		[32, 80, 14, 14]	(102,900)	False	
		└─Sequential (5)			[32, 80,
14, 14]		[32, 112, 14, 14]	--	False	
		└─MBConv (0)			[32, 80,
14, 14]		[32, 112, 14, 14]	(126,004)	False	
		└─MBConv (1)			[32, 112,
14, 14]		[32, 112, 14, 14]	(208,572)	False	
		└─MBConv (2)			[32, 112,
14, 14]		[32, 112, 14, 14]	(208,572)	False	
		└─Sequential (6)			[32, 112,
14, 14]		[32, 192, 7, 7]	--	False	
		└─MBConv (0)			[32, 112,
14, 14]		[32, 192, 7, 7]	(262,492)	False	
		└─MBConv (1)			[32, 192,
7, 7]		[32, 192, 7, 7]	(587,952)	False	
		└─MBConv (2)			[32, 192,
7, 7]		[32, 192, 7, 7]	(587,952)	False	
		└─MBConv (3)			[32, 192,
7, 7]		[32, 192, 7, 7]	(587,952)	False	
		└─Sequential (7)			[32, 192,
7, 7]		[32, 320, 7, 7]	--	False	
		└─MBConv (0)			[32, 192,
7, 7]		[32, 320, 7, 7]	(717,232)	False	
		└─Conv2dNormActivation (8)			[32, 320,
7, 7]		[32, 1280, 7, 7]	--	False	
		└─Conv2d (0)			[32, 320,
7, 7]		[32, 1280, 7, 7]	(409,600)	False	
		└─BatchNorm2d (1)			[32,
1280, 7, 7]		[32, 1280, 7, 7]	(2,560)	False	
		└─SiLU (2)			[32,
1280, 7, 7]		[32, 1280, 7, 7]	--	--	
		└─AdaptiveAvgPool2d (avgpool)			[32,
1280, 7, 7]		[32, 1280, 1, 1]	--	--	
		└─Sequential (classifier)			[32,
1280]		[32, 3]	--	True	
		└─Dropout (0)			[32,
1280]		[32, 1280]	--	--	
		└─Linear (1)			[32,
1280]		[32, 3]	3,843	True	
=====					
=====					
Total params: 4,011,391					
Trainable params: 3,843					
Non-trainable params: 4,007,548					
Total mult-adds (G): 12.31					
=====					
=====					
Input size (MB): 19.27					


```
Forward/backward pass size (MB): 3452.09
Params size (MB): 16.05
Estimated Total Size (MB): 3487.41
```

4. Train a single model and track results

```
# Define loss function optimizer
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

To track experiments, we're going to use TensorBoard:

<https://www.tensorflow.org/tensorboard/>

And to interact with TensorBoard, we can use PyTorch's SummaryWriter -

<https://pytorch.org/docs/stable/tensorboard.html>

- Also see here:
<https://pytorch.org/docs/stable/tensorboard.html#torch.utils.tensorboard.writer.SummaryWriter>

```
# Setup a SummaryWriter
from torch.utils.tensorboard import SummaryWriter
writer = SummaryWriter()
writer

<torch.utils.tensorboard.writer.SummaryWriter at 0x7fe2d94bb5d0>

from tqdm.auto import tqdm
from typing import Dict, List, Tuple

from going_modular.going_modular.engine import train_step, test_step

def train(model: torch.nn.Module,
          train_dataloader: torch.utils.data.DataLoader,
          test_dataloader: torch.utils.data.DataLoader,
          optimizer: torch.optim.Optimizer,
          loss_fn: torch.nn.Module,
          epochs: int,
          device: torch.device) -> Dict[str, List]:
    """Trains and tests a PyTorch model.

    Passes a target PyTorch models through train_step() and
    test_step()
    functions for a number of epochs, training and testing the model
    in the same epoch loop.

    Calculates, prints and stores evaluation metrics throughout.

    Args:
```

model: A PyTorch model to be trained and tested.
train_dataloader: A DataLoader instance for the model to be trained on.
test_dataloader: A DataLoader instance for the model to be tested on.
optimizer: A PyTorch optimizer to help minimize the loss function.
loss_fn: A PyTorch loss function to calculate loss on both datasets.
epochs: An integer indicating how many epochs to train for.
device: A target device to compute on (e.g. "cuda" or "cpu").

Returns:

A dictionary of training and testing loss as well as training and testing accuracy metrics. Each metric has a value in a list for each epoch.

In the form: {train_loss: [...],
 train_acc: [...],
 test_loss: [...],
 test_acc: [...]}

For example if training for epochs=2:
 {train_loss: [2.0616, 1.0537],
 train_acc: [0.3945, 0.3945],
 test_loss: [1.2641, 1.5706],
 test_acc: [0.3400, 0.2973]}

"""

Create empty results dictionary

```
results = {"train_loss": [],  
          "train_acc": [],  
          "test_loss": [],  
          "test_acc": []  
}
```

Loop through training and testing steps for a number of epochs

```
for epoch in tqdm(range(epochs)):  
    train_loss, train_acc = train_step(model=model,  
                                        dataloader=train_dataloader,  
                                        loss_fn=loss_fn,  
                                        optimizer=optimizer,  
                                        device=device)  
  
    test_loss, test_acc = test_step(model=model,  
                                    dataloader=test_dataloader,  
                                    loss_fn=loss_fn,  
                                    device=device)
```

Print out what's happening

```
print(  
    f"Epoch: {epoch+1} | "  
    f"train_loss: {train_loss:.4f} | "  
    f"train_acc: {train_acc:.4f} | "  
    f"test_loss: {test_loss:.4f} | "
```

```

        f"test_acc: {test_acc:.4f}"
    )

    # Update results dictionary
    results["train_loss"].append(train_loss)
    results["train_acc"].append(train_acc)
    results["test_loss"].append(test_loss)
    results["test_acc"].append(test_acc)

    ### New: Experiment tracking ###
    # See SummaryWriter documentation
    writer.add_scalars(main_tag="Loss",
                       tag_scalar_dict={"train_loss": train_loss,
                                         "test_loss": test_loss},
                       global_step=epoch)

    writer.add_scalars(main_tag="Accuracy",
                       tag_scalar_dict={"train_acc": train_acc,
                                         "test_acc": test_acc},
                       global_step=epoch)

    writer.add_graph(model=model,
                     input_to_model=torch.randn(32, 3, 224,
224).to(device))

    # Close the writer
    writer.close()
    ### End new ###

    # Return the filled results at the end of the epochs
    return results

# Train model
# Note: not using engine.train(), since we updated the train()
function above
set_seeds()
results = train(model=model,
                 train_dataloader=train_dataloader,
                 test_dataloader=test_dataloader,
                 optimizer=optimizer,
                 loss_fn=loss_fn,
                 epochs=5,
                 device=device)

{"model_id": "3799609be39c4acab9659d16719049d3", "version_major": 2, "version_minor": 0}

Epoch: 1 | train_loss: 1.0929 | train_acc: 0.4023 | test_loss: 0.9125
| test_acc: 0.5502
Epoch: 2 | train_loss: 0.8966 | train_acc: 0.6562 | test_loss: 0.7839
| test_acc: 0.8561

```

```
Epoch: 3 | train_loss: 0.8045 | train_acc: 0.7422 | test_loss: 0.6716  
| test_acc: 0.8864  
Epoch: 4 | train_loss: 0.6787 | train_acc: 0.7305 | test_loss: 0.6697  
| test_acc: 0.8258  
Epoch: 5 | train_loss: 0.7066 | train_acc: 0.7188 | test_loss: 0.6737  
| test_acc: 0.7737
```

results

```
{'test_acc': [0.5501893939393939,  
0.8560606060606061,  
0.8863636363636364,  
0.8257575757575758,  
0.7736742424242425],  
'test_loss': [0.9124558766682943,  
0.7839208642641703,  
0.671596626440684,  
0.6696672836939493,  
0.6737416386604309],  
'train_acc': [0.40234375, 0.65625, 0.7421875, 0.73046875, 0.71875],  
'train_loss': [1.0929433777928352,  
0.8965702056884766,  
0.8045311793684959,  
0.6786761954426765,  
0.706612229347229]}
```

5. View our model's results with TensorBoard

There are a few ways to view TensorBoard results, see them here:

https://www.learnpytorch.io/07_pytorch_experiment_tracking/#5-view-our-models-results-in-tensorboard

```
# Let's view our experiments from within the notebook  
%load_ext tensorboard  
%tensorboard --logdir runs  
  
<IPython.core.display.Javascript object>
```

6. Create a function to prepare a SummaryWriter() instance

By default our SummaryWriter() class saves to log_dir.

How about if we wanted to save different experiments to different folders?

In essence, one experiment = one folder.

For example, we'd like to track:

- Experiment date/timestamp
- Experiment name

- Model name
- Extra - is there anything else that should be tracked?

Let's create a function to create a `SummaryWriter()` instance to take all of these things into account.

So ideally we end up tracking experiments to a directory:

`runs/YYYY-MM-DD/experiment_name/model_name/extra`

```
from torch.utils.tensorboard import SummaryWriter
def create_writer(experiment_name: str,
                  model_name: str,
                  extra: str = None):
    """Creates a torch.utils.tensorboard.writer.SummaryWriter() instance
    tracking to a specific directory."""
    from datetime import datetime
    import os

    # Get timestamp of current date in reverse order
    timestamp = datetime.now().strftime("%Y-%m-%d")

    if extra:
        # Create log directory path
        log_dir = os.path.join("runs", timestamp, experiment_name,
                               model_name, extra)
    else:
        log_dir = os.path.join("runs", timestamp, experiment_name,
                               model_name)
    print(f"[INFO] Created SummaryWriter saving to {log_dir}")
    return SummaryWriter(log_dir=log_dir)

example_writer = create_writer(experiment_name="data_10_percent",
                               model_name="effnetb0",
                               extra="5_epochs")

example_writer

[INFO] Created SummaryWriter saving to
runs/2022-06-28/data_10_percent/effnetb0/5_epochs

<torch.utils.tensorboard.writer.SummaryWriter at 0x7fe1d40ff0d0>
```

6.1 Update the `train()` function to include a writer parameter

```
from tqdm.auto import tqdm
from typing import Dict, List, Tuple

from going_modular.going_modular.engine import train_step, test_step

def train(model: torch.nn.Module,
          train_data_loader: torch.utils.data.DataLoader,
```

```

        test_dataloader: torch.utils.data.DataLoader,
        optimizer: torch.optim.Optimizer,
        loss_fn: torch.nn.Module,
        epochs: int,
        device: torch.device,
        writer: torch.utils.tensorboard.writer.SummaryWriter) ->
Dict[str, List]:
    """Trains and tests a PyTorch model.

    Passes a target PyTorch models through train_step() and
    test_step()
    functions for a number of epochs, training and testing the model
    in the same epoch loop.

    Calculates, prints and stores evaluation metrics throughout.

    Args:
        model: A PyTorch model to be trained and tested.
        train_dataloader: A DataLoader instance for the model to be
        trained on.
        test_dataloader: A DataLoader instance for the model to be tested
        on.
        optimizer: A PyTorch optimizer to help minimize the loss function.
        loss_fn: A PyTorch loss function to calculate loss on both
        datasets.
        epochs: An integer indicating how many epochs to train for.
        device: A target device to compute on (e.g. "cuda" or "cpu").

    Returns:
        A dictionary of training and testing loss as well as training and
        testing accuracy metrics. Each metric has a value in a list for
        each epoch.
        In the form: {train_loss: [...],
                      train_acc: [...],
                      test_loss: [...],
                      test_acc: [...]}
        For example if training for epochs=2:
        {train_loss: [2.0616, 1.0537],
          train_acc: [0.3945, 0.3945],
          test_loss: [1.2641, 1.5706],
          test_acc: [0.3400, 0.2973]}
    """
    # Create empty results dictionary
    results = {"train_loss": [],
              "train_acc": [],
              "test_loss": [],
              "test_acc": []
             }

    # Loop through training and testing steps for a number of epochs

```

```

for epoch in tqdm(range(epochs)):
    train_loss, train_acc = train_step(model=model,
                                       dataloader=train_dataloader,
                                       loss_fn=loss_fn,
                                       optimizer=optimizer,
                                       device=device)

    test_loss, test_acc = test_step(model=model,
                                    dataloader=test_dataloader,
                                    loss_fn=loss_fn,
                                    device=device)

    # Print out what's happening
    print(
        f"Epoch: {epoch+1} | "
        f"train_loss: {train_loss:.4f} | "
        f"train_acc: {train_acc:.4f} | "
        f"test_loss: {test_loss:.4f} | "
        f"test_acc: {test_acc:.4f}"
    )

    # Update results dictionary
    results["train_loss"].append(train_loss)
    results["train_acc"].append(train_acc)
    results["test_loss"].append(test_loss)
    results["test_acc"].append(test_acc)

    ### New: Experiment tracking ###
    if writer:
        # See SummaryWriter documentation
        writer.add_scalars(main_tag="Loss",
                           tag_scalar_dict={"train_loss":
train_loss,
                                           "test_loss": test_loss},
                           global_step=epoch)

        writer.add_scalars(main_tag="Accuracy",
                           tag_scalar_dict={"train_acc": train_acc,
                                           "test_acc": test_acc},
                           global_step=epoch)

        writer.add_graph(model=model,
                           input_to_model=torch.randn(32, 3, 224,
224).to(device))

        # Close the writer
        writer.close()

    else:
        pass
    ### End new ###

```

```
# Return the filled results at the end of the epochs  
return results
```

7. Setting up a series of modelling experiments

- Challenge: Setup 2x modelling experiments with effnetb0, pizza, steak, sushi data and train one model for 5 epochs and another model for 10 epochs

7.1 What kind of experiments should you run?

The number of machine learning experiments you can run, is like the number of different models you can build... almost limitless.

However, you can't test everything...

So what should you test?

- Change the number of epochs
- Change the number of hidden layers/units
- Change the amount of data (right now we're using 10% of the Food101 dataset for pizza, steak, sushi)
- Change the learning rate
- Try different kinds of data augmentation
- Choose a different model architecture

This is why transfer learning is so powerful, because, it's a working model that you can apply to your own problem.

7.2 What experiments are we going to run?

We're going to turn three dials:

1. Model size - EffnetB0 vs EffnetB2 (in terms of number of parameters)
2. Dataset size - 10% of pizza, steak, sushi images vs 20% (generally more data = better results)
3. Training time - 5 epochs vs 10 epochs (generally longer training time = better results, up to a point)

To begin, we're still keeping things relatively small so that our experiments run quickly.

Our goal: a model that is well performing but still small enough to run on a mobile device or web browser, so FoodVision Mini can come to life.

If you had infinite compute + time, you should basically always choose the biggest model and biggest dataset you can. See: <http://www.incompleteideas.net/Incldeas/BitterLesson.html>

7.3 Download different datasets

We want two datasets:

1. Pizza, steak, sushi 10% -
https://github.com/mrdbourke/pytorch-deep-learning/raw/main/data/pizza_steak_sushi.zip
2. Pizza, steak, sushi 20% -
https://github.com/mrdbourke/pytorch-deep-learning/raw/main/data/pizza_steak_sushi_20_percent.zip

They were created with:

https://github.com/mrdbourke/pytorch-deep-learning/blob/main/extras/04_custom_data_creation.ipynb

```
# Download 10 percent and 20 percent datasets
data_10_percent_path =
download_data(source="https://github.com/mrdbourke/pytorch-deep-
learning/raw/main/data/pizza_steak_sushi.zip",
              destination="pizza_steak_sushi")

data_20_percent_path =
download_data(source="https://github.com/mrdbourke/pytorch-deep-
learning/raw/main/data/pizza_steak_sushi_20_percent.zip",
              destination="pizza_steak_sushi_20_percent")

[INFO] data/pizza_steak_sushi directory already exists, skipping
download.
[INFO] Did not find data/pizza_steak_sushi_20_percent directory,
creating one...
[INFO] Downloading pizza_steak_sushi_20_percent.zip from
https://github.com/mrdbourke/pytorch-deep-learning/raw/main/data/pizza
_steak_sushi_20_percent.zip...
[INFO] Unzipping pizza_steak_sushi_20_percent.zip data...
```

7.4 Transform Datasets and Create DataLoaders

We'll transform our data in a few ways:

1. Resize the images to (224, 224)
2. Make sure image tensor values are between [0, 1]
3. Normalize the images so they have the same data distribution as ImageNet

```
# Setup training directory paths
train_dir_10_percent = data_10_percent_path / "train"
train_dir_20_percent = data_20_percent_path / "train"

# Setup the test directory
test_dir = data_10_percent_path / "test"

train_dir_10_percent, train_dir_20_percent, test_dir
```

```

(PosixPath('data/pizza_steak_sushi/train'),
 PosixPath('data/pizza_steak_sushi_20_percent/train'),
 PosixPath('data/pizza_steak_sushi/test'))

from torchvision import transforms

# Setup ImageNet normalization levels
# See here: https://pytorch.org/vision/0.12/models.html
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])

# Compose transforms into a pipeline
simple_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    normalize
])

BATCH_SIZE = 32

# Create 10% training and test DataLoaders
train_dataloader_10_percent, test_dataloader, class_names =
data_setup.create_dataloaders(train_dir=train_dir_10_percent,

test_dir=test_dir,

transform=simple_transform,

batch_size=BATCH_SIZE)

# Create 20% training and test DataLoaders
train_dataloader_20_percent, test_dataloader, class_names =
data_setup.create_dataloaders(train_dir=train_dir_20_percent,

test_dir=test_dir,

transform=simple_transform,

batch_size=BATCH_SIZE)

print(f"Number of batches of size {BATCH_SIZE} in 10% train data:
{len(train_dataloader_10_percent)}")
print(f"Number of batches of size {BATCH_SIZE} in 20% train data:
{len(train_dataloader_20_percent)}")
print(f"Number of batches of size {BATCH_SIZE} in 10% test data:
{len(test_dataloader)}")
print(f"Class names: {class_names}")

Number of batches of size 32 in 10% train data: 8
Number of batches of size 32 in 20% train data: 15

```

Number of batches of size 32 in 10% test data: 3
Class names: ['pizza', 'steak', 'sushi']

7.5 Create feature extractor models

We want two functions:

1. Creates a `torchvision.models.efficientnet_b0()` feature extractor with a frozen backbone/base layers and a custom classifier head (EffNetB0).
2. Creates a `torchvision.models.efficientnet_b2()` feature extractor with a frozen backbone/base layers and a custom classifier head (EffNetB2).

```
import torchvision

# Create an EffNetB2
effnetb2_weights = torchvision.models.EfficientNet_B2_Weights.DEFAULT
# "DEFAULT" = best available
effnetb2 =
torchvision.models.efficientnet_b2(weights=effnetb2_weights)

# effnetb2

Downloading:
"https://download.pytorch.org/models/efficientnet_b2_rwightman-
bcd3f34b7.pth" to
/root/.cache/torch/hub/checkpoints/efficientnet_b2_rwightman-
bcd3f34b7.pth

{"model_id": "321952803de84cecaa56a0b250abefa0", "version_major": 2, "version_minor": 0}

summary(model=effnetb2,
        input_size=(32, 3, 224, 224),
        verbose=0,
        col_names=["input_size", "output_size", "num_params",
"trainable"],
        col_width=20,
        row_settings=["var_names"])
```

```
=====
=====
Layer (type (var_name))          Input
Shape      Output Shape      Param #    Trainable
=====
=====
EfficientNet (EfficientNet)      [32, 3,
224, 224]    [32, 1000]    --         True
├─Sequential (features)         [32, 3,
224, 224]    [32, 1408, 7, 7]  --         True
|   └─Conv2dNormActivation (0)  [32, 3,
```

224, 224]	[32, 32, 112, 112]	--	True	
	└─Conv2d (0)			[32, 3,
224, 224]	[32, 32, 112, 112]	864	True	
	└─BatchNorm2d (1)			[32, 32,
112, 112]	[32, 32, 112, 112]	64	True	
	└─SiLU (2)			[32, 32,
112, 112]	[32, 32, 112, 112]	--	--	
	└─Sequential (1)			[32, 32,
112, 112]	[32, 16, 112, 112]	--	True	
	└─MBConv (0)			[32, 32,
112, 112]	[32, 16, 112, 112]	1,448	True	
	└─MBConv (1)			[32, 16,
112, 112]	[32, 16, 112, 112]	612	True	
	└─Sequential (2)			[32, 16,
112, 112]	[32, 24, 56, 56]	--	True	
	└─MBConv (0)			[32, 16,
112, 112]	[32, 24, 56, 56]	6,004	True	
	└─MBConv (1)			[32, 24,
56, 56]	[32, 24, 56, 56]	10,710	True	
	└─MBConv (2)			[32, 24,
56, 56]	[32, 24, 56, 56]	10,710	True	
	└─Sequential (3)			[32, 24,
56, 56]	[32, 48, 28, 28]	--	True	
	└─MBConv (0)			[32, 24,
56, 56]	[32, 48, 28, 28]	16,518	True	
	└─MBConv (1)			[32, 48,
28, 28]	[32, 48, 28, 28]	43,308	True	
	└─MBConv (2)			[32, 48,
28, 28]	[32, 48, 28, 28]	43,308	True	
	└─Sequential (4)			[32, 48,
28, 28]	[32, 88, 14, 14]	--	True	
	└─MBConv (0)			[32, 48,
28, 28]	[32, 88, 14, 14]	50,300	True	
	└─MBConv (1)			[32, 88,
14, 14]	[32, 88, 14, 14]	123,750	True	
	└─MBConv (2)			[32, 88,
14, 14]	[32, 88, 14, 14]	123,750	True	
	└─MBConv (3)			[32, 88,
14, 14]	[32, 88, 14, 14]	123,750	True	
	└─Sequential (5)			[32, 88,
14, 14]	[32, 120, 14, 14]	--	True	
	└─MBConv (0)			[32, 88,
14, 14]	[32, 120, 14, 14]	149,158	True	
	└─MBConv (1)			[32, 120,
14, 14]	[32, 120, 14, 14]	237,870	True	
	└─MBConv (2)			[32, 120,
14, 14]	[32, 120, 14, 14]	237,870	True	
	└─MBConv (3)			[32, 120,
14, 14]	[32, 120, 14, 14]	237,870	True	

	└Sequential (6)				[32, 120,
14, 14]	[32, 208, 7, 7]	--		True	
	└└MBConv (0)				[32, 120,
14, 14]	[32, 208, 7, 7]	301,406		True	
	└└MBConv (1)				[32, 208,
7, 7]	[32, 208, 7, 7]	686,868		True	
	└└MBConv (2)				[32, 208,
7, 7]	[32, 208, 7, 7]	686,868		True	
	└└MBConv (3)				[32, 208,
7, 7]	[32, 208, 7, 7]	686,868		True	
	└└MBConv (4)				[32, 208,
7, 7]	[32, 208, 7, 7]	686,868		True	
	└Sequential (7)				[32, 208,
7, 7]	[32, 352, 7, 7]	--		True	
	└└MBConv (0)				[32, 208,
7, 7]	[32, 352, 7, 7]	846,900		True	
	└└MBConv (1)				[32, 352,
7, 7]	[32, 352, 7, 7]	1,888,920		True	
	└└Conv2dNormActivation (8)				[32, 352,
7, 7]	[32, 1408, 7, 7]	--		True	
	└└└Conv2d (0)				[32, 352,
7, 7]	[32, 1408, 7, 7]	495,616		True	
	└└└BatchNorm2d (1)				[32,
1408, 7, 7]	[32, 1408, 7, 7]	2,816		True	
	└└└SiLU (2)				[32,
1408, 7, 7]	[32, 1408, 7, 7]	--		--	
	└AdaptiveAvgPool2d (avgpool)				[32,
1408, 7, 7]	[32, 1408, 1, 1]	--		--	
	└Sequential (classifier)				[32,
1408]	[32, 1000]	--		True	
	└└Dropout (0)				[32,
1408]	[32, 1408]	--		--	
	└└Linear (1)				[32,
1408]	[32, 1000]	1,409,000		True	
=====					
Total params: 9,109,994					
Trainable params: 9,109,994					
Non-trainable params: 0					
Total mult-adds (G): 21.09					
=====					
Input size (MB): 19.27					
Forward/backward pass size (MB): 5017.79					
Params size (MB): 36.44					
Estimated Total Size (MB): 5073.49					
=====					

```

import torchvision
from torch import nn

OUT_FEATURES = len(class_names)

# Create an EffNetB0 feature extractor
def create_effnetb0():
    # Get the weights and setup a model
    weights = torchvision.models.EfficientNet_B0_Weights.DEFAULT
    model =
    torchvision.models.efficientnet_b0(weights=weights).to(device)

    # Freeze the base model layers
    for param in model.features.parameters():
        param.requires_grad = False

    # Change the classifier head
    set_seeds()
    model.classifier = nn.Sequential(
        nn.Dropout(p=0.2, inplace=True),
        nn.Linear(in_features=1280, out_features=OUT_FEATURES)
    ).to(device)

    # Give the model a name
    model.name = "effnetb0"
    print(f"[INFO] Created new {model.name} model...")
    return model

# Create an EffNetB2 feature extractor
def create_effnetb2():
    # Get the weights and setup a model
    weights = torchvision.models.EfficientNet_B2_Weights.DEFAULT
    model =
    torchvision.models.efficientnet_b2(weights=weights).to(device)

    # Freeze the base model layers
    for param in model.features.parameters():
        param.requires_grad = False

    # Change the classifier head
    set_seeds()
    model.classifier = nn.Sequential(
        nn.Dropout(p=0.3, inplace=True),
        nn.Linear(in_features=1408, out_features=OUT_FEATURES)
    ).to(device)

    # Give the model a name
    model.name = "effnetb2"
    print(f"[INFO] Created new {model.name} model...")
    return model

```

```

effnetb2.classifier
Sequential(
  (0): Dropout(p=0.3, inplace=True)
  (1): Linear(in_features=1408, out_features=1000, bias=True)
)

created_model_test_effnetb2 = create_effnetb2()
created_model_test_effnetb0 = create_effnetb0()

[INFO] Created new effnetb2 model...
[INFO] Created new effnetb0 model...

# Check out EffNetB2 feature extractor
summary(model=created_model_test_effnetb2,
        input_size=(32, 3, 224, 224),
        verbose=0,
        col_names=["input_size", "output_size", "num_params",
"trainable"],
        col_width=20,
        row_settings=["var_names"])

```

=====			
=====			
Layer (type (var_name))	Output Shape	Param #	Input Trainable
=====			
=====			
EfficientNet (EfficientNet)			[32, 3,
224, 224]	[32, 3]	--	Partial
└─Sequential (features)			[32, 3,
224, 224]	[32, 1408, 7, 7]	--	False
└─Conv2dNormActivation (0)			[32, 3,
224, 224]	[32, 32, 112, 112]	--	False
└─Conv2d (0)			[32, 3,
224, 224]	[32, 32, 112, 112]	(864)	False
└─BatchNorm2d (1)			[32, 32,
112, 112]	[32, 32, 112, 112]	(64)	False
└─SiLU (2)			[32, 32,
112, 112]	[32, 32, 112, 112]	--	--
└─Sequential (1)			[32, 32,
112, 112]	[32, 16, 112, 112]	--	False
└─MBConv (0)			[32, 32,
112, 112]	[32, 16, 112, 112]	(1,448)	False
└─MBConv (1)			[32, 16,
112, 112]	[32, 16, 112, 112]	(612)	False
└─Sequential (2)			[32, 16,
112, 112]	[32, 24, 56, 56]	--	False
└─MBConv (0)			[32, 16,
112, 112]	[32, 24, 56, 56]	(6,004)	False

		└─MBConv (1)				[32, 24,
56, 56]		[32, 24, 56, 56]	(10,710)	False		
		└─MBConv (2)				[32, 24,
56, 56]		[32, 24, 56, 56]	(10,710)	False		
		└─Sequential (3)				[32, 24,
56, 56]		[32, 48, 28, 28]	--	False		
		└─MBConv (0)				[32, 24,
56, 56]		[32, 48, 28, 28]	(16,518)	False		
		└─MBConv (1)				[32, 48,
28, 28]		[32, 48, 28, 28]	(43,308)	False		
		└─MBConv (2)				[32, 48,
28, 28]		[32, 48, 28, 28]	(43,308)	False		
		└─Sequential (4)				[32, 48,
28, 28]		[32, 88, 14, 14]	--	False		
		└─MBConv (0)				[32, 48,
28, 28]		[32, 88, 14, 14]	(50,300)	False		
		└─MBConv (1)				[32, 88,
14, 14]		[32, 88, 14, 14]	(123,750)	False		
		└─MBConv (2)				[32, 88,
14, 14]		[32, 88, 14, 14]	(123,750)	False		
		└─MBConv (3)				[32, 88,
14, 14]		[32, 88, 14, 14]	(123,750)	False		
		└─Sequential (5)				[32, 88,
14, 14]		[32, 120, 14, 14]	--	False		
		└─MBConv (0)				[32, 88,
14, 14]		[32, 120, 14, 14]	(149,158)	False		
		└─MBConv (1)				[32, 120,
14, 14]		[32, 120, 14, 14]	(237,870)	False		
		└─MBConv (2)				[32, 120,
14, 14]		[32, 120, 14, 14]	(237,870)	False		
		└─MBConv (3)				[32, 120,
14, 14]		[32, 120, 14, 14]	(237,870)	False		
		└─Sequential (6)				[32, 120,
14, 14]		[32, 208, 7, 7]	--	False		
		└─MBConv (0)				[32, 120,
14, 14]		[32, 208, 7, 7]	(301,406)	False		
		└─MBConv (1)				[32, 208,
7, 7]		[32, 208, 7, 7]	(686,868)	False		
		└─MBConv (2)				[32, 208,
7, 7]		[32, 208, 7, 7]	(686,868)	False		
		└─MBConv (3)				[32, 208,
7, 7]		[32, 208, 7, 7]	(686,868)	False		
		└─MBConv (4)				[32, 208,
7, 7]		[32, 208, 7, 7]	(686,868)	False		
		└─Sequential (7)				[32, 208,
7, 7]		[32, 352, 7, 7]	--	False		
		└─MBConv (0)				[32, 208,
7, 7]		[32, 352, 7, 7]	(846,900)	False		
		└─MBConv (1)				[32, 352,

7, 7]	[32, 352, 7, 7]	(1,888,920)	False	
	└─Conv2dNormActivation (8)			[32, 352,
7, 7]	[32, 1408, 7, 7]	--	False	
	└─Conv2d (0)			[32, 352,
7, 7]	[32, 1408, 7, 7]	(495,616)	False	
	└─BatchNorm2d (1)			[32,
1408, 7, 7]	[32, 1408, 7, 7]	(2,816)	False	
	└─SiLU (2)			[32,
1408, 7, 7]	[32, 1408, 7, 7]	--	--	
	└─AdaptiveAvgPool2d (avgpool)			[32,
1408, 7, 7]	[32, 1408, 1, 1]	--	--	
	└─Sequential (classifier)			[32,
1408]	[32, 3]	--	True	
	└─Dropout (0)			[32,
1408]	[32, 1408]	--	--	
	└─Linear (1)			[32,
1408]	[32, 3]	4,227	True	

```

=====
Total params: 7,705,221
Trainable params: 4,227
Non-trainable params: 7,700,994
Total mult-adds (G): 21.04
=====

```

```

=====
Input size (MB): 19.27
Forward/backward pass size (MB): 5017.53
Params size (MB): 30.82
Estimated Total Size (MB): 5067.62
=====

```

Check out EffNetB0 feature extractor model

```

summary(model=created_model_test_effnetb0,
        input_size=(32, 3, 224, 224),
        verbose=0,
        col_names=["input_size", "output_size", "num_params",
"trainable"],
        col_width=20,
        row_settings=["var_names"])

```

Layer (type (var_name))	Output Shape	Param #	Input Trainable
EfficientNet (EfficientNet)	[32, 3,		
224, 224]	[32, 3]	--	Partial
└─Sequential (features)			[32, 3,

224, 224]	[32, 1280, 7, 7]	--	False	
	└─Conv2dNormActivation (0)			[32, 3,
224, 224]	[32, 32, 112, 112]	--	False	
	└─Conv2d (0)			[32, 3,
224, 224]	[32, 32, 112, 112]	(864)	False	
	└─BatchNorm2d (1)			[32, 32,
112, 112]	[32, 32, 112, 112]	(64)	False	
	└─SiLU (2)			[32, 32,
112, 112]	[32, 32, 112, 112]	--	--	
	└─Sequential (1)			[32, 32,
112, 112]	[32, 16, 112, 112]	--	False	
	└─MBConv (0)			[32, 32,
112, 112]	[32, 16, 112, 112]	(1,448)	False	
	└─Sequential (2)			[32, 16,
112, 112]	[32, 24, 56, 56]	--	False	
	└─MBConv (0)			[32, 16,
112, 112]	[32, 24, 56, 56]	(6,004)	False	
	└─MBConv (1)			[32, 24,
56, 56]	[32, 24, 56, 56]	(10,710)	False	
	└─Sequential (3)			[32, 24,
56, 56]	[32, 40, 28, 28]	--	False	
	└─MBConv (0)			[32, 24,
56, 56]	[32, 40, 28, 28]	(15,350)	False	
	└─MBConv (1)			[32, 40,
28, 28]	[32, 40, 28, 28]	(31,290)	False	
	└─Sequential (4)			[32, 40,
28, 28]	[32, 80, 14, 14]	--	False	
	└─MBConv (0)			[32, 40,
28, 28]	[32, 80, 14, 14]	(37,130)	False	
	└─MBConv (1)			[32, 80,
14, 14]	[32, 80, 14, 14]	(102,900)	False	
	└─MBConv (2)			[32, 80,
14, 14]	[32, 80, 14, 14]	(102,900)	False	
	└─Sequential (5)			[32, 80,
14, 14]	[32, 112, 14, 14]	--	False	
	└─MBConv (0)			[32, 80,
14, 14]	[32, 112, 14, 14]	(126,004)	False	
	└─MBConv (1)			[32, 112,
14, 14]	[32, 112, 14, 14]	(208,572)	False	
	└─MBConv (2)			[32, 112,
14, 14]	[32, 112, 14, 14]	(208,572)	False	
	└─Sequential (6)			[32, 112,
14, 14]	[32, 192, 7, 7]	--	False	
	└─MBConv (0)			[32, 112,
14, 14]	[32, 192, 7, 7]	(262,492)	False	
	└─MBConv (1)			[32, 192,
7, 7]	[32, 192, 7, 7]	(587,952)	False	
	└─MBConv (2)			[32, 192,
7, 7]	[32, 192, 7, 7]	(587,952)	False	

		└─MBConv (3)			[32, 192,
7, 7]		[32, 192, 7, 7]	(587,952)	False	
		└─Sequential (7)			[32, 192,
7, 7]		[32, 320, 7, 7]	--	False	
		└─MBConv (0)			[32, 192,
7, 7]		[32, 320, 7, 7]	(717,232)	False	
		└─Conv2dNormActivation (8)			[32, 320,
7, 7]		[32, 1280, 7, 7]	--	False	
		└─Conv2d (0)			[32, 320,
7, 7]		[32, 1280, 7, 7]	(409,600)	False	
		└─BatchNorm2d (1)			[32,
1280, 7, 7]		[32, 1280, 7, 7]	(2,560)	False	
		└─SiLU (2)			[32,
1280, 7, 7]		[32, 1280, 7, 7]	--	--	
		└─AdaptiveAvgPool2d (avgpool)			[32,
1280, 7, 7]		[32, 1280, 1, 1]	--	--	
		└─Sequential (classifier)			[32,
1280]		[32, 3]	--	True	
		└─Dropout (0)			[32,
1280]		[32, 1280]	--	--	
		└─Linear (1)			[32,
1280]		[32, 3]	3,843	True	

=====

Total params: 4,011,391
Trainable params: 3,843
Non-trainable params: 4,007,548
Total mult-adds (G): 12.31

=====

Input size (MB): 19.27
Forward/backward pass size (MB): 3452.09
Params size (MB): 16.05
Estimated Total Size (MB): 3487.41

=====

7.6 Create experiments and set up training code

```
# Create epoch list
num_epochs = [5, 10]

# Create models list (need to create a new model for each experiment)
models = ["effnetb0", "effnetb2"]

# Create a DataLoaders dictionary
train_dataloaders = {"data_10_percent": train_dataloader_10_percent,
                     "data_20_percent": train_dataloader_20_percent}
```

```

%%time
from going_modular.going_modular.utils import save_model

# Set seeds
set_seeds(seed=42)

# Keep track of experiment numbers
experiment_number = 0

# Loop through each DataLoader
for dataloader_name, train_dataloader in train_dataloaders.items():
    # Loop through the epochs
    for epochs in num_epochs:
        # Loop through each model name and create a new model instance
        for model_name in models:

            # Print out info
            experiment_number += 1
            print(f"[INFO] Experiment number: {experiment_number}")
            print(f"[INFO] Model: {model_name}")
            print(f"[INFO] DataLoader: {dataloader_name}")
            print(f"[INFO] Number of epochs: {epochs}")

            # Select and create the model
            if model_name == "effnetb0":
                model = create_effnetb0()
            else:
                model = create_effnetb2()

            # Create a new loss and optimizer for every model
            loss_fn = nn.CrossEntropyLoss()
            optimizer = torch.optim.Adam(params=model.parameters(),
lr=0.001)

            # Train target model with target dataloader and track
            experiments
            # Note: using train() rather than engine.train()
            train(model=model,
                train_dataloader=train_dataloader,
                test_dataloader=test_dataloader,
                optimizer=optimizer,
                loss_fn=loss_fn,
                epochs=epochs,
                device=device,
                writer=create_writer(experiment_name=dataloader_name,
                    model_name=model_name,
                    extra=f"{epochs}_epochs"))

            # Save the model to file so we can import it later if need be
            save_filepath =

```

```
f"07_{model_name}_{dataloader_name}_{epochs}_epochs.pth"
    save_model(model=model,
                target_dir="models",
                model_name=save_filepath)
    print("-"*50 + "\n")
```

```
[INFO] Experiment number: 1
[INFO] Model: effnetb0
[INFO] DataLoader: data_10_percent
[INFO] Number of epochs: 5
[INFO] Created new effnetb0 model...
[INFO] Created SummaryWriter saving to
runs/2022-06-28/data_10_percent/effnetb0/5_epochs
```

```
{"model_id": "44f62855a2fe49d98c4f219e379ff8b8", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 1.0433 | train_acc: 0.4805 | test_loss: 0.9283
| test_acc: 0.4782
Epoch: 2 | train_loss: 0.9400 | train_acc: 0.5469 | test_loss: 0.8346
| test_acc: 0.5492
Epoch: 3 | train_loss: 0.8243 | train_acc: 0.6953 | test_loss: 0.7293
| test_acc: 0.8248
Epoch: 4 | train_loss: 0.7029 | train_acc: 0.7773 | test_loss: 0.6177
| test_acc: 0.8759
Epoch: 5 | train_loss: 0.6150 | train_acc: 0.8867 | test_loss: 0.5830
| test_acc: 0.8864
[INFO] Saving model to:
models/07_effnetb0_data_10_percent_5_epochs.pth
-----
```

```
[INFO] Experiment number: 2
[INFO] Model: effnetb2
[INFO] DataLoader: data_10_percent
[INFO] Number of epochs: 5
[INFO] Created new effnetb2 model...
[INFO] Created SummaryWriter saving to
runs/2022-06-28/data_10_percent/effnetb2/5_epochs
```

```
{"model_id": "1fa379b53b1b4bdab7864f77d09ec4a1", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 1.0861 | train_acc: 0.4023 | test_loss: 0.9619
| test_acc: 0.6903
Epoch: 2 | train_loss: 0.8960 | train_acc: 0.6250 | test_loss: 0.9031
| test_acc: 0.6818
Epoch: 3 | train_loss: 0.8458 | train_acc: 0.6641 | test_loss: 0.8036
| test_acc: 0.8049
Epoch: 4 | train_loss: 0.7027 | train_acc: 0.8516 | test_loss: 0.6922
| test_acc: 0.9176
```

```
Epoch: 5 | train_loss: 0.7005 | train_acc: 0.7422 | test_loss: 0.6509  
| test_acc: 0.8968  
[INFO] Saving model to:  
models/07_effnetb2_data_10_percent_5_epochs.pth  
-----
```

```
[INFO] Experiment number: 3  
[INFO] Model: effnetb0  
[INFO] DataLoader: data_10_percent  
[INFO] Number of epochs: 10  
[INFO] Created new effnetb0 model...  
[INFO] Created SummaryWriter saving to  
runs/2022-06-28/data_10_percent/effnetb0/10_epochs
```

```
{"model_id": "6cd298c9251543bb9546efffb441841", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 1.0433 | train_acc: 0.4805 | test_loss: 0.9283  
| test_acc: 0.4782  
Epoch: 2 | train_loss: 0.9400 | train_acc: 0.5469 | test_loss: 0.8346  
| test_acc: 0.5492  
Epoch: 3 | train_loss: 0.8243 | train_acc: 0.6953 | test_loss: 0.7293  
| test_acc: 0.8248  
Epoch: 4 | train_loss: 0.7029 | train_acc: 0.7773 | test_loss: 0.6177  
| test_acc: 0.8759  
Epoch: 5 | train_loss: 0.6150 | train_acc: 0.8867 | test_loss: 0.5830  
| test_acc: 0.8864  
Epoch: 6 | train_loss: 0.5431 | train_acc: 0.8750 | test_loss: 0.5943  
| test_acc: 0.8561  
Epoch: 7 | train_loss: 0.6385 | train_acc: 0.7109 | test_loss: 0.5933  
| test_acc: 0.8561  
Epoch: 8 | train_loss: 0.5075 | train_acc: 0.8008 | test_loss: 0.5382  
| test_acc: 0.8968  
Epoch: 9 | train_loss: 0.4688 | train_acc: 0.9180 | test_loss: 0.5318  
| test_acc: 0.8759  
Epoch: 10 | train_loss: 0.5795 | train_acc: 0.7344 | test_loss: 0.4907  
| test_acc: 0.8759  
[INFO] Saving model to:  
models/07_effnetb0_data_10_percent_10_epochs.pth  
-----
```

```
[INFO] Experiment number: 4  
[INFO] Model: effnetb2  
[INFO] DataLoader: data_10_percent  
[INFO] Number of epochs: 10  
[INFO] Created new effnetb2 model...  
[INFO] Created SummaryWriter saving to  
runs/2022-06-28/data_10_percent/effnetb2/10_epochs
```

```
{"model_id": "c2c40c14fe8b4a579ba9352af0764e63", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 1.0861 | train_acc: 0.4023 | test_loss: 0.9619  
| test_acc: 0.6903  
Epoch: 2 | train_loss: 0.8960 | train_acc: 0.6250 | test_loss: 0.9031  
| test_acc: 0.6818  
Epoch: 3 | train_loss: 0.8458 | train_acc: 0.6641 | test_loss: 0.8036  
| test_acc: 0.8049  
Epoch: 4 | train_loss: 0.7027 | train_acc: 0.8516 | test_loss: 0.6922  
| test_acc: 0.9176  
Epoch: 5 | train_loss: 0.7005 | train_acc: 0.7422 | test_loss: 0.6509  
| test_acc: 0.8968  
Epoch: 6 | train_loss: 0.5988 | train_acc: 0.8984 | test_loss: 0.6597  
| test_acc: 0.8769  
Epoch: 7 | train_loss: 0.6206 | train_acc: 0.8047 | test_loss: 0.6037  
| test_acc: 0.9384  
Epoch: 8 | train_loss: 0.5341 | train_acc: 0.8164 | test_loss: 0.6002  
| test_acc: 0.8570  
Epoch: 9 | train_loss: 0.5031 | train_acc: 0.9258 | test_loss: 0.5606  
| test_acc: 0.8570  
Epoch: 10 | train_loss: 0.5309 | train_acc: 0.8008 | test_loss: 0.5462  
| test_acc: 0.8977
```

```
[INFO] Saving model to:  
models/07_effnetb2_data_10_percent_10_epochs.pth
```

```
-----  
[INFO] Experiment number: 5  
[INFO] Model: effnetb0  
[INFO] DataLoader: data_20_percent  
[INFO] Number of epochs: 5  
[INFO] Created new effnetb0 model...  
[INFO] Created SummaryWriter saving to  
runs/2022-06-28/data_20_percent/effnetb0/5_epochs
```

```
{"model_id": "4ac449dde12d4269aee5ddc1ae103fd3", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 0.9645 | train_acc: 0.5938 | test_loss: 0.6673  
| test_acc: 0.8864  
Epoch: 2 | train_loss: 0.6941 | train_acc: 0.7833 | test_loss: 0.5834  
| test_acc: 0.9072  
Epoch: 3 | train_loss: 0.5817 | train_acc: 0.8313 | test_loss: 0.5024  
| test_acc: 0.9280  
Epoch: 4 | train_loss: 0.4582 | train_acc: 0.8917 | test_loss: 0.4186  
| test_acc: 0.9072  
Epoch: 5 | train_loss: 0.4519 | train_acc: 0.8604 | test_loss: 0.3827  
| test_acc: 0.9280
```

```
[INFO] Saving model to:  
models/07_effnetb0_data_20_percent_5_epochs.pth
```

```
-----  
[INFO] Experiment number: 6  
[INFO] Model: effnetb2  
[INFO] DataLoader: data_20_percent  
[INFO] Number of epochs: 5  
[INFO] Created new effnetb2 model...  
[INFO] Created SummaryWriter saving to  
runs/2022-06-28/data_20_percent/effnetb2/5_epochs
```

```
{"model_id": "ac3e6e1aead3481886228b86cf25df01", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 0.9883 | train_acc: 0.5271 | test_loss: 0.7823  
| test_acc: 0.8049  
Epoch: 2 | train_loss: 0.7125 | train_acc: 0.8250 | test_loss: 0.6505  
| test_acc: 0.8864  
Epoch: 3 | train_loss: 0.5772 | train_acc: 0.8917 | test_loss: 0.5523  
| test_acc: 0.9384  
Epoch: 4 | train_loss: 0.5209 | train_acc: 0.8583 | test_loss: 0.5109  
| test_acc: 0.9280  
Epoch: 5 | train_loss: 0.4635 | train_acc: 0.8667 | test_loss: 0.4445  
| test_acc: 0.9384  
[INFO] Saving model to:  
models/07_effnetb2_data_20_percent_5_epochs.pth  
-----
```

```
[INFO] Experiment number: 7  
[INFO] Model: effnetb0  
[INFO] DataLoader: data_20_percent  
[INFO] Number of epochs: 10  
[INFO] Created new effnetb0 model...  
[INFO] Created SummaryWriter saving to  
runs/2022-06-28/data_20_percent/effnetb0/10_epochs
```

```
{"model_id": "91695204baa94fb3bec484d4dcca8228", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 0.9645 | train_acc: 0.5938 | test_loss: 0.6673  
| test_acc: 0.8864  
Epoch: 2 | train_loss: 0.6941 | train_acc: 0.7833 | test_loss: 0.5834  
| test_acc: 0.9072  
Epoch: 3 | train_loss: 0.5817 | train_acc: 0.8313 | test_loss: 0.5024  
| test_acc: 0.9280  
Epoch: 4 | train_loss: 0.4582 | train_acc: 0.8917 | test_loss: 0.4186  
| test_acc: 0.9072  
Epoch: 5 | train_loss: 0.4519 | train_acc: 0.8604 | test_loss: 0.3827  
| test_acc: 0.9280  
Epoch: 6 | train_loss: 0.4218 | train_acc: 0.8917 | test_loss: 0.3734  
| test_acc: 0.9384
```



```
Epoch: 7 | train_loss: 0.4005 | train_acc: 0.8646 | test_loss: 0.3395  
| test_acc: 0.9176  
Epoch: 8 | train_loss: 0.3453 | train_acc: 0.8896 | test_loss: 0.3497  
| test_acc: 0.9384  
Epoch: 9 | train_loss: 0.3740 | train_acc: 0.8938 | test_loss: 0.3457  
| test_acc: 0.9384  
Epoch: 10 | train_loss: 0.3839 | train_acc: 0.8938 | test_loss: 0.2856  
| test_acc: 0.9176  
[INFO] Saving model to:  
models/07_effnetb0_data_20_percent_10_epochs.pth  
-----
```

```
[INFO] Experiment number: 8  
[INFO] Model: effnetb2  
[INFO] DataLoader: data_20_percent  
[INFO] Number of epochs: 10  
[INFO] Created new effnetb2 model...  
[INFO] Created SummaryWriter saving to  
runs/2022-06-28/data_20_percent/effnetb2/10_epochs
```

```
{"model_id": "847bcefa15744539b3f41d4658ba6cd0", "version_major": 2, "version_minor": 0}
```

```
Epoch: 1 | train_loss: 0.9883 | train_acc: 0.5271 | test_loss: 0.7823  
| test_acc: 0.8049  
Epoch: 2 | train_loss: 0.7125 | train_acc: 0.8250 | test_loss: 0.6505  
| test_acc: 0.8864  
Epoch: 3 | train_loss: 0.5772 | train_acc: 0.8917 | test_loss: 0.5523  
| test_acc: 0.9384  
Epoch: 4 | train_loss: 0.5209 | train_acc: 0.8583 | test_loss: 0.5109  
| test_acc: 0.9280  
Epoch: 5 | train_loss: 0.4635 | train_acc: 0.8667 | test_loss: 0.4445  
| test_acc: 0.9384  
Epoch: 6 | train_loss: 0.3801 | train_acc: 0.9146 | test_loss: 0.4582  
| test_acc: 0.9280  
Epoch: 7 | train_loss: 0.3597 | train_acc: 0.8938 | test_loss: 0.4195  
| test_acc: 0.9280  
Epoch: 8 | train_loss: 0.3280 | train_acc: 0.9021 | test_loss: 0.4135  
| test_acc: 0.9280  
Epoch: 9 | train_loss: 0.3467 | train_acc: 0.8896 | test_loss: 0.4325  
| test_acc: 0.8674  
Epoch: 10 | train_loss: 0.3743 | train_acc: 0.8583 | test_loss: 0.3768  
| test_acc: 0.9489  
[INFO] Saving model to:  
models/07_effnetb2_data_20_percent_10_epochs.pth  
-----
```

```
CPU times: user 2min 51s, sys: 27.2 s, total: 3min 18s  
Wall time: 5min 29s
```

8. View experiments in TensorBoard

We've *experiment, experiment, experiment!*

Now let's *visualize, visualize, visualize!*

```
# Let's view our experiments within TensorBoard from within the notebook
```

```
%load_ext tensorboard
```

```
%tensorboard --logdir runs
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

```
Reusing TensorBoard on port 6006 (pid 375), started 0:23:43 ago. (Use '!kill 375' to kill it.)
```

```
<IPython.core.display.Javascript object>
```

The best performing model was:

- Model: EffNetB2
- Dataset: Pizza, steak, sushi 20%
- Epochs: 10

And the overall trend of all the results was that more data, bigger model and longer training time generally led to better results.

```
# # Upload the results to TensorBoard.dev (uncomment to try it out)
```

```
!tensorboard dev upload --logdir runs \
```

```
    --name "07. PyTorch Experiment Tracking: FoodVision Mini model  
result (video)" \
```

```
    --description "Comparing results of different model size, training  
data amount and training time."
```

```
    --one_shot
```

You can view the experiments publically at TensorBoard.dev:

<https://tensorboard.dev/experiment/4yXqVFNyQQymt4yfvvs6sA/>

9. Load in the best model and make predictions with it

This is our best model filepath:

```
models/07_effnetb2_data_20_percent_10_epochs.pth
```

```
# Setup best model filepath
```

```
best_model_path = "models/07_effnetb2_data_20_percent_10_epochs.pth"
```

```
# Instantiate a new instance of EffNetB2 (to load in the saved  
state_dict())
```

```

best_model = create_effnetb2()

# Load the saved best model state_dict()
best_model.load_state_dict(torch.load(best_model_path))

[INFO] Created new effnetb2 model...

<All keys matched successfully>

```

Our goal: create a FoodVision Mini model that performs well enough and is able to run on a mobile device/web browser.

```

# Check the model file size
from pathlib import Path

# Get the model size in bytes then convert it to megabytes
effnetb2_model_size = Path(best_model_path).stat().st_size //
(1024*1024)
print(f"EfficientNetB2 feature extractor model size:
{effnetb2_model_size} MB")

EfficientNetB2 feature extractor model size: 29 MB

# Import function to make prediction on images and plot them
from going_modular.going_modular.predictions import
pred_and_plot_image

# Get a random list of 3 image path names from the test dataset
import random
num_images_to_plot = 3
test_image_path_list = list(Path(data_20_percent_path /
"test").glob("*/*.jpg"))
test_image_path_sample = random.sample(test_image_path_list,
k=num_images_to_plot)

for image_path in test_image_path_sample:
    pred_and_plot_image(model=best_model,
                        image_path=image_path,
                        class_names=class_names,
                        image_size=(224, 224))

```

Pred: pizza | Prob: 0.961



Pred: pizza | Prob: 0.927



Pred: sushi | Prob: 0.656



9.1 Predict on a custom image with the best model

```
# Download custom image
import requests
from pathlib import Path

# Setup custom image path
custom_image_path = Path("data/04-pizza-dad.jpeg")

# Download the image if it doesn't already exist
if not custom_image_path.is_file():
    with open(custom_image_path, "wb") as f:
        # When downloading from GitHub, need to use the "raw" file
        link
        request =
requests.get("https://raw.githubusercontent.com/mrdbourke/pytorch-
deep-learning/main/images/04-pizza-dad.jpeg")
        print(f"Downloading {custom_image_path}...")
        f.write(request.content)
else:
    print(f"{custom_image_path} already exists, skipping download.")

Downloading data/04-pizza-dad.jpeg...

# Predict on our own custom image
pred_and_plot_image(model=model,
                    image_path=custom_image_path,
                    class_names=class_names)
```

Pred: pizza | Prob: 0.982

